MDPI

*Review*

# Blockchain-Assisted Mission Trust in Air–Surface–Underwater Unmanned Systems: A Review and Conceptual Framework

**Anonymous Author(s)**

Author affiliations withheld for double-anonymous review
* Correspondence details withheld for double-anonymous review

## Abstract

Air–surface–underwater missions combine heterogeneous vehicles, intermittent links, and administrative domains. A relying decision must connect physical observations, protected records, shared authority, and fitness for use. This structured narrative review introduces the Object–Assurance–Lifecycle (OAL) framework and applies it to 23 full-text mechanism studies through 49 claim records, a coverage heatmap, and scoped evidence profiles. The coded sample concentrates on admission integrity and data provenance/governance, while transitions to disconnected mission use require complementary evidence. A typed dependency graph and context–temporal–graph loop organize reassessment when observations, keys, policies, or dependencies change. A partition–revocation–reconciliation case specifies coalition roles, a Fabric-oriented application pattern, constrained verification, and an acoustic-link payload budget. An executable synthetic replay checks cryptographic bindings and temporal/authorization rules; a policy-validating witnessed log produces the same decisions under matched rules. Ledger value is therefore conditional on joint-governance requirements and implementation cost. Requirement-based consensus selection, falsifiable propositions, and mappings to existing assurance-case and maritime frameworks define an evaluation agenda. The contribution is a reproducible analytical synthesis and design illustration; comparative deployment performance and integrated mission effectiveness remain empirical questions.

**Keywords:** mission trust; unmanned systems; UAV–USV–UUV cooperation; blockchain; data provenance; remote attestation; zero trust; maritime autonomy

## 1. Introduction

Air–surface–underwater missions combine unmanned aerial vehicles (UAVs), unmanned surface vehicles (USVs), and unmanned underwater vehicles (UUVs) for observation, inspection, search, and environmental monitoring. Their complementary mobility and sensing capabilities create dependencies across radio and acoustic links, gateways, processing services, and administrative owners [1–3]. UAV and marine-system reviews identify mobility, limited energy, intermittent communication, localization uncertainty, and difficult recovery as operational constraints [4–7]. A mission decision must combine evidence generated under these different conditions.

This creates a *mission-trust gap*. Authentication establishes a principal's identity; a content hash binds bytes; replicated consensus establishes agreement under a fault model. A relying decision additionally requires evidence that the observation is physically credible, sufficiently current, and suitable for its intended consequence. A signed sonar report can be inaccurate, a genuine image can be stale, and a correctly ordered record can originate from

Received:
Accepted:
Published:

**Copyright:** © 2026 by the authors. Submitted to *Blockchains* for possible open access publication under the terms and conditions of the Creative Commons Attribution (CC BY) license.

a captured device. UAV and underwater-security research describes the corresponding spoofing, jamming, capture, routing, and data-manipulation threats [8–12]. Connecting these controls to the eventual use of a product requires explicit claim semantics.

Blockchain proposals address swarm authentication, data acquisition, membership, provenance, and audit [13–16]. Reviews of UAV and maritime applications describe coordination and sharing opportunities [17–21]. Their value depends on which parties must jointly govern state, which faults must be tolerated, and whether an authenticated database or independently witnessed log already satisfies those requirements. The blockchain oracle literature names the external-input problem directly: agreement about recorded assertions does not establish their correspondence with physical events [22,23].

Adjacent mechanisms supply complementary evidence. Zero trust provides resource-specific authorization [24,25]; remote attestation separates device evidence, appraisal, and relying decisions [26]; provenance records producers and transformations [27,28]; and trust inference models reputation, uncertainty, and changing behavior [29–31]. These mechanisms use different vocabularies and evaluate different claims. Their composition is the central subject of this review.

The proposed *Object–Assurance–Lifecycle* (OAL) framework relates four objects—entity, platform/device, link/path, and data product—to four assurance layers—observation validity, commitment and provenance integrity, ledger and governance consistency, and mission-use suitability—across six lifecycle stages. A typed dependency graph and context–temporal–graph (CTG) loop preserve shared causes, event time, and uncertainty when applying this structure to a decision. OAL supplies the claim organization; established inference, provenance, policy, and consensus mechanisms supply implementations.

The review makes three connected contributions. First, it applies OAL to a disclosed full-text sample of 23 mechanism studies, producing a reproducible coverage map with claim locators and evidence profiles. Second, a partition–revocation–reconciliation case connects the framework to a specified permissioned-ledger pattern, an executable policy replay, and a communication budget. A policy-validating witnessed-log baseline exposes where the two architectures provide equivalent decisions. Third, requirement-based consensus selection and falsifiable research propositions translate the synthesis into an evaluation agenda. The analytical contribution concerns evidence organization and conditional technology selection; the replay evaluates specified rules rather than an integrated vehicle deployment.

The paper proceeds from review methods and operational context to the literature synthesis, OAL mapping, dependency reasoning, ledger design, evaluation practice, and research agenda. Civil missions and publicly documented coalition settings provide the application context. Limitations of coverage, implementation evidence, and transfer to operational missions are consolidated in Section 10.

## 2. Review Scope and Methods

### *2.1. Identification and Selection*

This is a structured narrative review. The original bibliography contained 141 references spanning unmanned cooperation, communication and security, trust inference, provenance, attestation, and distributed governance. The retained discovery script queried OpenAlex on 6 August 2026 using 32 English search strings, a publication-date filter starting on 1 January 2016, and the 25 highest-ranked results per query. Deduplication by DOI or OpenAlex identifier produced 571 candidate records. Queries included “blockchain UAV swarm”, “underwater wireless sensor network trust management”, “UAV USV UUV cooperative systems”, “remote attestation IoT survey RATS”, and “blockchain network partition finality”. The complete query list and retrieval settings accompany the supple-

mentary materials. Citation following and targeted retrieval added earlier foundations and official specifications. Crossref metadata supported bibliographic checking; it was not a second independent screening database.

Selection favored English sources that described cross-domain operating constraints, specified a security or trust mechanism applicable to a mission boundary, or supplied a relevant standard, dataset, or simulator. Duplicate records and sources without a substantive connection to those questions were excluded from the synthesis. Reviews supplied context and citation leads; original mechanism studies supplied the analytical sample. Foundational work predating 2016 was retained for reputation, distributed consistency, provenance, and feedback architectures. During revision, targeted searches of publisher sites, author repositories, IETF, NIST, IMO, DNV, and standards organizations added the bodies of work identified by the reviewers and checked developments through 22 October 2026.

The historical workflow did not retain a complete record of title/abstract rejection decisions. Accordingly, the 571 candidates and 141 original citations describe discovery and the original bibliography, respectively; they are not a reconstructed systematic-review screening flow. The revision supplies a disposition for every original citation and separates full-text analytical coding from contextual use. Fourteen additional original mechanism candidates remain outside the numerical map because their full-text coding was not completed. This explicit boundary permits inspection of coverage without inferring evidence from titles or abstracts.

### *2.2. Full-Text Coding and Reproducible Synthesis*

The analytical sample comprises 23 original mechanism contributions for which full texts were inspected: four UAV authentication/acquisition studies, two underwater trust studies, and 17 transferable studies of reputation, provenance, attestation, access control, shared state, or observation appraisal. These are purposively selected studies within the narrative review, rather than a census of the field. Control/planning papers, reviews, standards, and datasets remain contextual or evaluation resources. Source identifiers S01–S23, extracted claims, supporting sections/pages, deployment scope, and coding decisions are available in `oal-study-mapping.csv`. The full-text PDFs are not redistributed.

The coding unit is an explicitly described mechanism claim. Each record associates an object with an assurance layer and one or more supported lifecycle stages. Several claims can be coded for one paper; unrelated object, layer, and stage labels are never expanded into an automatic Cartesian product. The map contains 49 claim records. A heatmap cell counts distinct studies associated with that exact object–layer–stage combination; a study is counted once per cell. Counts therefore describe coverage in the coded sample and are not additive sample sizes or measures of effectiveness.

Evidence profiles record conceptual argument (L0), analytical/security reasoning (L1), simulation or recorded-data evaluation (L2), implemented prototype/testbed (L3), physical operating trials (L4), and sustained multi-party operation (L5). These categories are complementary evidence forms, not an ordinal quality score. A profile is attached to the stated mechanism and documented evaluation scope; an implementation benchmark does not validate every assurance claim made in its surrounding architecture. Recorded field data analyzed offline are coded L2; a hardware prototype without described operating conditions is coded L3. Ambiguous deployment claims are assigned the narrower evidenced category and explained in the extraction note.

Coding was prepared as a single extraction pass and checked against the cited passages and aggregation rules; independent duplicate coding and inter-rater agreement were not performed. The released CSV, codebook, and figure-generation script make the clas-

sifications inspectable and revisable. Empty cells mean that no corresponding claim was observed in this sample. They do not establish the absence of relevant work in the wider literature.

**Table 1.** Scope and evidence roles in the revised synthesis.

| Role | Included material | Use in the synthesis |
|---|---|---|
| Operational context | UAV, USV, UUV, gateways, shore services; cooperative control and communication | Mission dependencies and deployment assumptions |
| Analytical sample | 23 full-text mechanism studies; 49 extracted claim records | OAL coverage, lifecycle association and scoped L0–L5 profiles |
| Contextual foundations | Reviews, reputation and distributed-system foundations, standards and governance cases | Vocabulary, alternatives and design requirements |
| Evaluation resources | Datasets, simulators and benchmark practices | Candidate tests; no inferred integrated mission validation |
| Revision case | Executable policy replay and exact application-byte accounting | Rule consistency and illustrative design trade-offs |

## 3. Operational Context and Trust Foundations

### *3.1. A Heterogeneous Mission System of Systems*

UAVs contribute rapid repositioning, line-of-sight radio relays, and electro-optical sensing. USVs carry larger energy stores and antennas and bridge acoustic and radio networks. UUVs operate close to underwater targets but often rely on low-rate acoustic links or delayed surfacing for large transfers. Communication surveys describe the resulting trade-offs among range, bandwidth, latency, energy, and alignment [32–35]. Topology, clock quality, and reachability can change within a mission.

Consider a civil coastal-response coalition comprising authority A, research operator B, communications provider C, and contracted data provider D. B operates a UUV and UAV; C supplies the USV relay; A approves mission authority; D contributes independent data and audit services. Each organization retains its administrative identity. A local operator can authenticate a device while lacking authority to change the coalition's membership, evidence policy, or revocation history. The case in Sections 6.4 and 7.3 uses this coalition to specify which changes require joint approval.

A USV may simultaneously relay acoustic traffic, cache data, provide time, and host fusion. Failure at that shared dependency affects several apparently independent products. End-to-end source signatures can preserve origin integrity through an untrusted relay, while selective withholding and common clock or transformation errors remain relevant. Figure 1 separates the constrained evidence path from shore-based shared-state services.

Raw imagery, sonar, detailed telemetry, and attestation evidence remain encrypted off chain. Shared records contain commitments, membership, policy/model versions, revocations, and audit judgments. Delay-tolerant forwarding can use Bundle Protocol version 7 and BPSec integrity/confidentiality services [36,37]. Store-and-forward delivery preserves messages across contacts; authority expiry and evidence-age policy are separate application rules. A successfully delivered old command is not thereby timely.

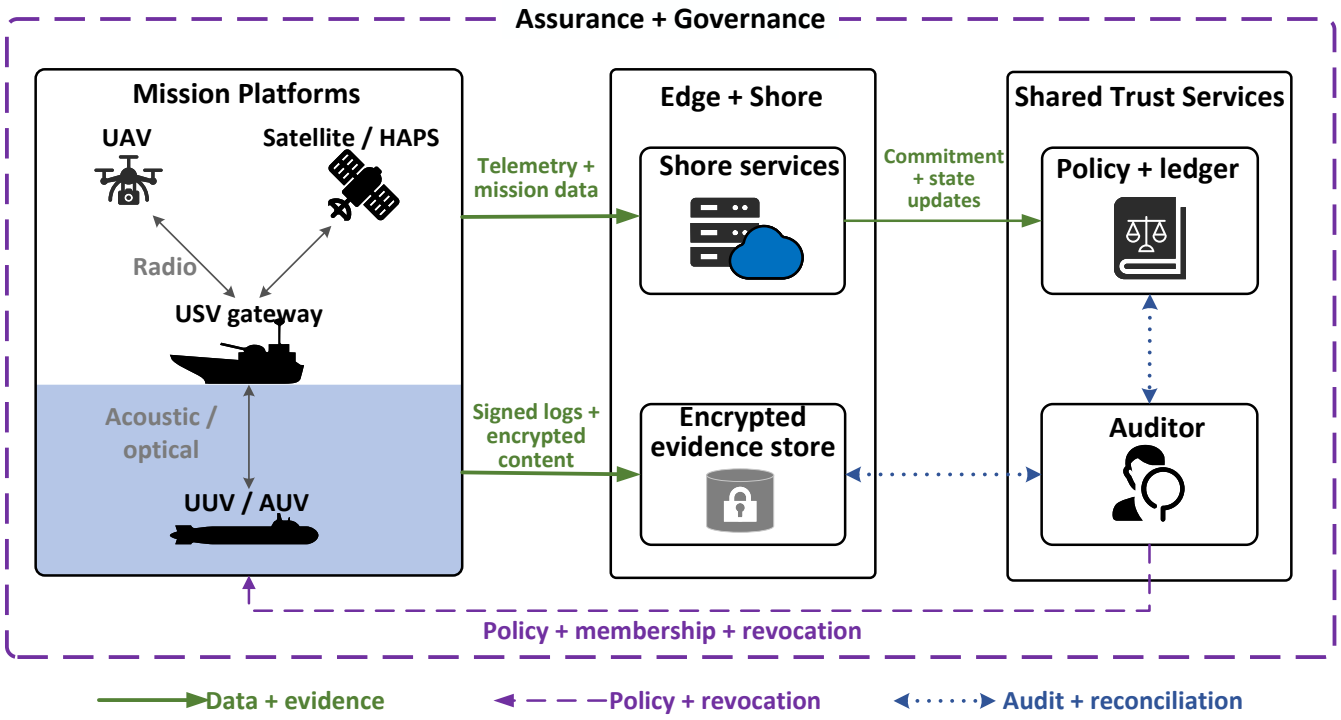



**Figure 1.** Cross-domain reference architecture. Shore organizations host peers and ordering nodes; the USV is a forwarding/cache dependency. Local sensing and safety continue within delegated authority during a gateway partition.

### 3.2. Mission Trust and the Relying Decision

*Mission trust* is a context-specific decision to rely on an object for a stated purpose, given available evidence and the consequences of error. *Trustworthiness* denotes the supporting properties: identity, measured endpoint state, link availability, provenance, freshness, observation quality, and accountable governance. The same image may support broad situational awareness while being unsuitable for collision avoidance.

For object $o$, time $t$, context $c$, and mission use $m$, appraisal and decision have distinct interfaces:

$$\begin{aligned} T_o(t,c,m) &= \mathcal{A}(o,\mathcal{E}_{\le t},c,m;\theta), \\ D_{r,o}(t,c,m) &= \pi_r(T_o(t,c,m),\mathcal{C}_t,m). \end{aligned} \tag{1}$$

Here $\mathcal{E}_{\le t}$ contains evidence received by decision time, including observation times and lineage. The appraisal $\mathcal{A}$ produces claim-specific outcomes and uncertainty under assumptions $\theta$. Relying party $r$ applies policy $\pi_r$ under authority, deadline, consequence, and resource constraints $\mathcal{C}_t$. The result may permit, restrict, defer, or reject use. Subjective opinions, intervals, ordinal states, and reputation values can implement this interface [31,38,39]. A probabilistic interpretation requires a defined event and calibration evidence.

A *trusted data environment* is the socio-technical setting that produces, transforms, shares, and governs this evidence. It contains principals, devices, paths, artifacts, and accountability processes. Zero trust contributes resource-centered policy and enforcement [24]; OAL records the additional observation, lineage, timing, and consequence claims relevant to mission reliance.

### 3.3. Four Complementary Assurance Claims

Observation validity concerns correspondence with the physical world. Commitment and provenance integrity concerns protected bytes and recorded origins or transformations.

Ledger and governance consistency concerns authorized shared-state transitions. Mission-use suitability concerns the particular decision and consequence. Figure 2 shows the distinct questions and supporting evidence.

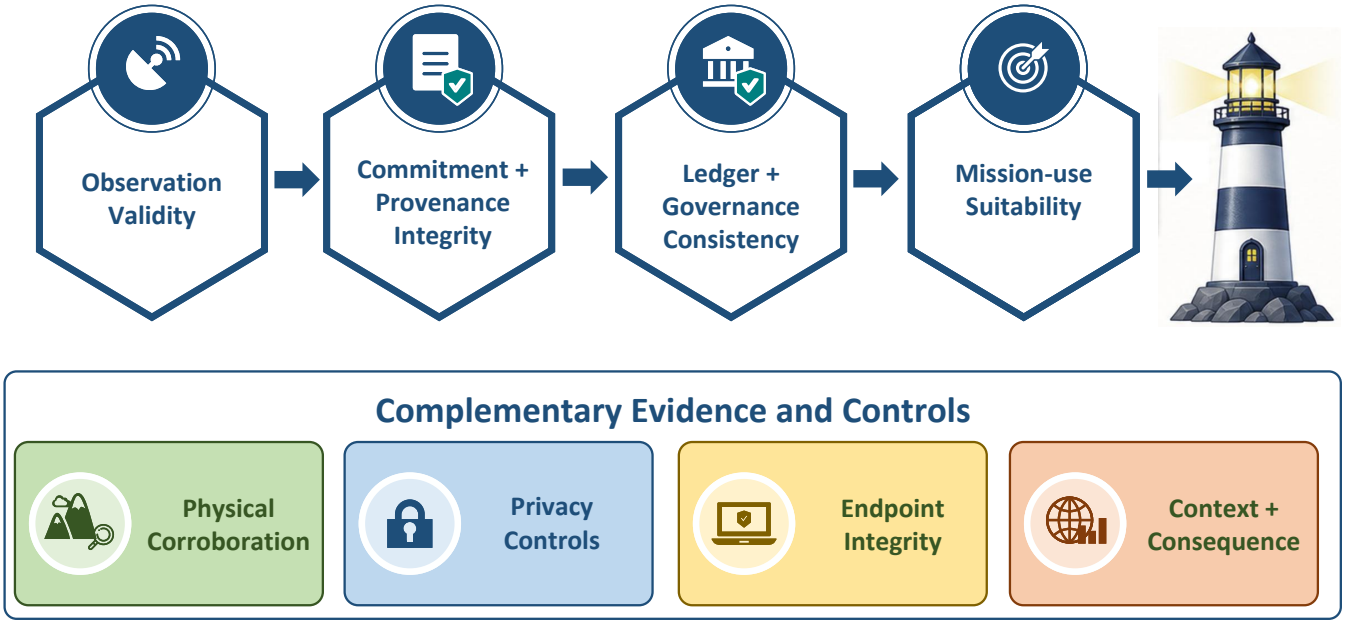



**Figure 2.** Four assurance claims and their evidence. The relation is an argument structure; an implementation may evaluate the checks in a different order.

A signature can remain valid after key theft; an accurate observation can become stale; and a consistent record can contain an erroneous sensor assertion. RATS appraisals support scoped measured-state claims [26], while physical diagnostics and corroboration support observation claims. The relying policy composes the applicable results, with explicit responses to failed or unknown requirements.

## 4. Complementary Research Landscapes

The literature connects three kinds of work: operational cooperation, mechanisms that assess or protect evidence, and services that coordinate shared records. Cooperative navigation and formation studies establish interactions among heterogeneous vehicles [40–43]. Their use of *consensus* concerns coordination of physical states. Ledger consensus concerns agreement on a replicated record under a fault model. OAL relates their outputs through identified observations, transformations, authorities, and intended uses.

### *4.1. Security and Evidence Appraisal*

Security surveys organize spoofing, jamming, capture, malicious routing, malware, and false-data injection by target and attack mechanism [8,11,44,45]. Underwater fading, noise, Doppler, and topology change can also cause loss or disagreement [33,46,47]. A trust estimate therefore needs a named property and an uncertainty status: an anomaly may justify restricted use while its cause remains unresolved.

Decision-tree and hidden-Markov approaches illustrate underwater trust inference [48,49]. The former evaluates simulated network behavior with an initial attack-free training period; the latter includes a numerical illustration with ten generated observations. Their evidence addresses the specified inference settings, rather than integrated coalition operation. Reputation-based peer selection [39] similarly supplies an appraisal mechanism whose meaning depends on the service and observed behavior. Survey taxonomies

distinguish direct observations, recommendations, context, and attacks on reputation [29,30,50,51].

Observation appraisal supplies a complementary strand. Signal-feature classification, network-assisted detection, and recorded-data validation have been studied for GNSS spoofing [52–54]. Multi-view opinion aggregation combines uncertain observations under declared assumptions [55]. These methods address evidence about the physical world. Their outputs become inputs to a relying policy with a specified consequence and validity interval.

### *4.2. Provenance, Attestation, and Access*

W3C PROV models entities, activities, and agents [27,28]. An OAL entity is an acting principal and usually maps to a PROV agent; an OAL data product maps to a PROV entity. Decision provenance extends the analysis along the data flows that lead to organizational or automated decisions [56]. It supports accountability when event capture and retained records cover the relevant transformations.

Attestation supplies measured endpoint-state evidence. Architectural proposals discuss attestation of ledger participants and autonomous IoT security [57,58]; implemented work evaluates attested channels and dual attestation of embedded devices and deployed ML artifacts [59,60]. RATS separates evidence, verifier appraisal, and the relying decision [26]. Zero trust and device baselines provide authorization and endpoint-capability requirements [24,25,61]. These boundaries matter during disconnection, when cached results and delegated permissions age independently.

### *4.3. Shared Records and Governance*

The coded UAV studies combine authenticated identities or data acquisition with ledger state [13–16]. Fabric provides an implemented execute–order–validate architecture [62]; IoT platforms apply it to sensing-record integrity, off-chain provenance, and access policy [63–65]. Food, health, and digital-twin studies extend these patterns to forecasting, federated learning, and lifecycle coordination [66–68]. Their recorded mechanisms and evaluation settings are summarized in Table 2 and mapped in Section 5.

**Table 2.** The 23 coded studies and their principal evidence scope. Identifiers match the supplementary claim records. Profiles describe reported evidence forms, not integrated maritime validation.

| Identifiers | Mechanism focus | Profiles | Evaluation scope |
|---|---|---|---|
| S01–S04 | UAV identity/authentication/acquisition | L1, L2, L3 | Security analysis, simulation and implemented components; BUS includes UAV and IoT hardware |
| S05–S07 | Underwater trust and P2P reputation | L1/L2 | Network or numerical evaluation and peer-selection simulation |
| S08–S10 | Fabric, sensing records, Hadoop provenance | L3 | Server and embedded-device prototypes; record processing and retrieval |
| S11–S14 | Attestation and autonomous IoT security | L0 or L3 | Two architectural contributions; two implemented attestation studies |
| S15–S17 | Food/health provenance and access control | L2 and/or L3 | Forecasting, FL applications and controlled access-policy requests |
| S18–S19 | Decision provenance and digital twins | L0 | Conceptual data-flow and lifecycle architectures |
| S20–S23 | GNSS and multi-view observation appraisal | L1/L2 | Signal/learning experiments, recorded data and aggregation analysis |

The oracle problem remains relevant across these applications: a ledger preserves the accepted assertion and its provenance, while external validity depends on the data source

and appraisal [22,23]. Transparency logs offer a natural alternative. Certificate Transparency specifies inclusion and consistency proofs, with external monitoring or witnessing needed to address inconsistent views [69,70]. Sigstore combines signing with transparency for software artifacts [71]. A witnessed log can also carry jointly authorized state transitions. The distinction to evaluate is therefore the required validation and governance service, together with its implementation cost, rather than the presence of hash chaining alone.

## 5. The OAL Framework and Literature Mapping

OAL classifies *what is judged*, *which claim is sought*, and *when that claim is needed*. A mechanism can support several cells; some combinations are inapplicable to a mission. Dependencies between claims are represented separately so that classification does not imply evidential independence.

### 5.1. Objects, Assurance Layers, and Lifecycle

An *entity* is an accountable principal: an organization, operator, service, or autonomous agent with delegated authority. A *platform/device* is the endpoint executing a function, including vehicles, gateways, sensors, and hosted compute instances. A vehicle can be considered an entity when exercising delegated authority and a platform when its firmware is measured. A *link/path* comprises physical media, intermediate nodes, security associations, and changing quality along an end-to-end route. A *data product* is an identified version of an observation, track, map, model, command, policy, or audit record. Corrections and transformations create new version identifiers.

**Table 3.** OAL objects and claim-specific evidence.

| Object | Evidence | Relying question |
|---|---|---|
| Entity (E) | Credential, role, delegation, accountable actions | Who acts and with which authority? |
| Platform (P) | Boot/runtime measurements, configuration, self-test, energy | Which endpoint executes, and in what measured state? |
| Link/path (L) | Authentication, route, delay, loss, custody and medium state | What protection and timeliness held along the path? |
| Data product (D) | Signature, content ID, times, lineage, uncertainty, residuals | Which version is this, and what supports its intended use? |

The assurance axis uses O (observation validity), C (commitment and provenance integrity), G (ledger and governance consistency), and M (mission-use suitability). O uses physical bounds, calibration, diagnostics, and corroboration. C binds producer, version, and transformation assertions through signatures and provenance. G addresses jointly authorized membership, order, policy, and recovery state. M applies purpose, freshness, consequence, uncertainty, and authority requirements. When a shared-state service is not needed, ledger-specific claims are inapplicable; equivalent governance claims may be met by another service.

Policies, models, and reference values are versioned data products; their deployed instances reside on platforms and their approval belongs to entities. This treatment makes the assurance infrastructure subject to the same analysis as mission observations. For each applicable cell, a specification states the object/version, required property, evidence source, validity interval, acceptance rule, responsible party, and response to missing evidence.

### 5.2. Observed Coverage in the Coded Sample

Figure 3 applies these axes to the 23 studies described in Section 2.2. Counts are generated from the released claim records, with one count per study per exact cell. Admission-

**Table 4.** Lifecycle stages and relevant state transitions.

| Stage | Required assurance | Examples of retained state |
|---|---|---|
| Admission (A) | Authorized principal and acceptable endpoint | Role, delegation, attestation reference |
| Sensing (S) | Credible time, place and observation | Product ID, quality and acquisition commitment |
| Transmission (T) | Protected content and adequate path | Relay/custody event and batch root |
| Fusion/sharing (F) | Identified inputs, transformation and access | Lineage, model and policy version |
| Decision/action (D) | Fit for the specified consequence now | Decision rationale and responsible principal |
| Recovery/audit (R) | Scope of invalidation and restored authority | Revocation, dispute, reconciled history |

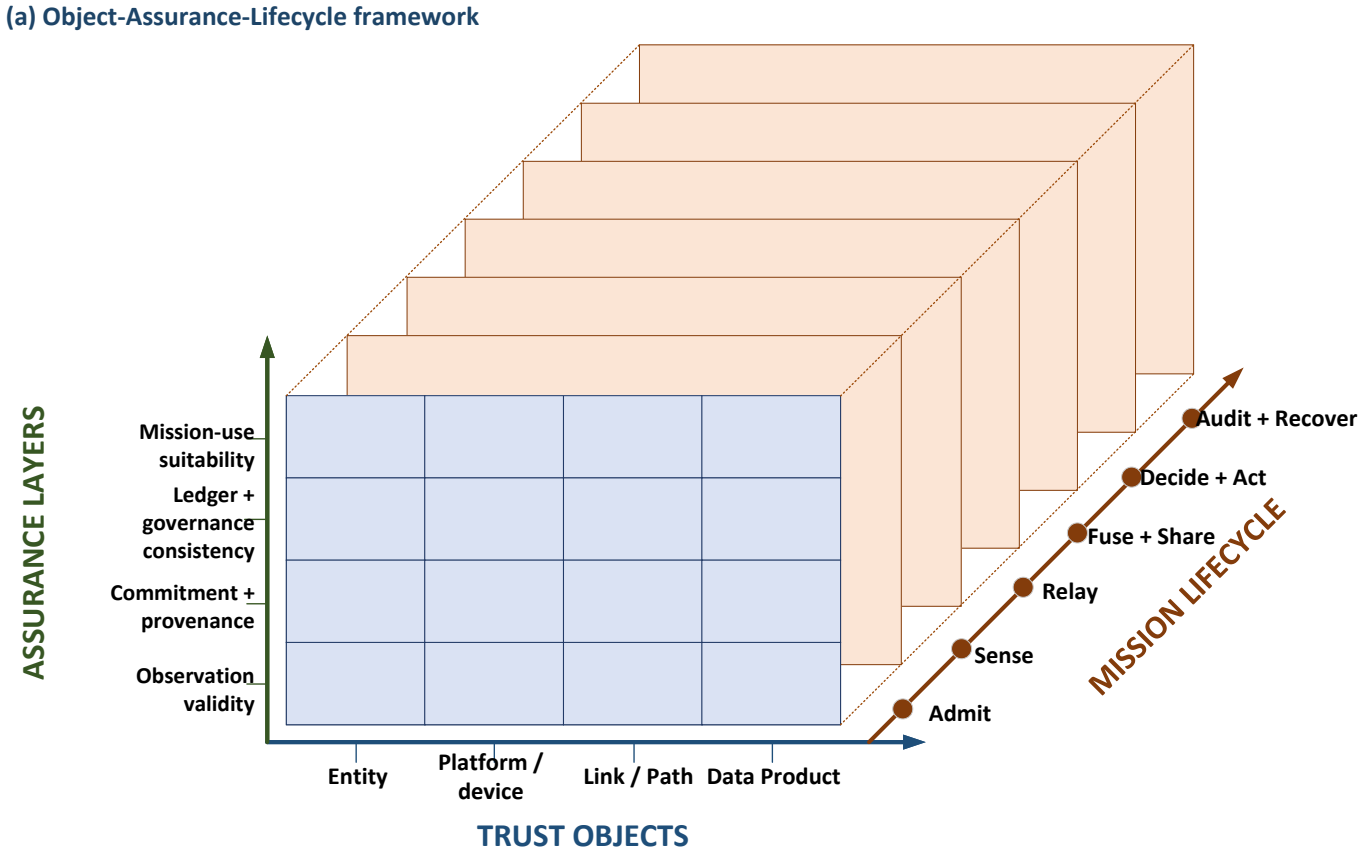
related identity and platform integrity, and data commitment/governance during fusion and audit, recur across the sample. Observation appraisal is concentrated in data products during sensing and fusion. The underwater studies address forwarding suitability using node or path evidence. These concentrations reflect the selected mechanisms and source availability.

The map makes three integration questions concrete. First, authentication at admission requires a continuity rule before later mission use: an old credential or attestation does not specify its validity during a partition. Second, provenance and shared-record integrity during fusion require complementary observation and consequence evidence. Third, recovery must identify affected descendants and distinguish event-time authority from later receipt of revocation. No coded claim establishes an integrated air–surface–underwater mission with all three transitions validated. That is a finding about the extracted evidence scope, not an assertion that cross-domain demonstrations do not exist.

Examples show the classification's practical effect. S02 maps UAV authentication and ledger registration to admission-related claims, while its cloud-peer implementation and simulated mission geometry remain visible in the evidence note. S10 maps Hadoop/Fabric provenance to data commitment and governance at fusion and recovery, without assigning sensor-truth validation. S12 supplies platform-attestation evidence at admission; its measured TLS configurations do not establish uninterrupted device integrity until a later action. S22 contributes recorded-data observation appraisal and is not recoded as a new integrated field trial. A design review can thus identify the extra evidence needed at each handover instead of transferring an entire paper's performance claim to a different object or stage.

The UAV subset illustrates why the evidence profile accompanies the cells. Han et al. implement a Fabric-based identity service on cloud servers and use a modeled UAV setting [13]. Tan et al. combine a lightweight authentication design with a Fabric implementation and simulated network geometry [14]. Karmakar et al. combine PUF-based authentication, clustering, and distributed ledger components [15]. BUS reports an implementation involving four IoT devices, three UAVs, a MEC server, and an Ethereum consortium [16]. Together these studies provide component precedents for admission and authenticated acquisition. The next mission-stage question is whether the accepted identity and observation remain appropriate after authority, endpoint state, or communication context changes.

The attestation subset separates architectural and implementation evidence. Hardjono and Smith analyze an attestation architecture for ledger participants, while ARCADIAN-IoT specifies a broader autonomous-security architecture [57,58]. Their coded profiles are conceptual. The implemented studies examine particular platform and protocol combinations




**(b) Coverage of 23 full-text studies / 49 claim records**

| | | Admission | Sensing | Relay | Fusion | Decision | Recovery |
|---|---|---|---|---|---|---|---|
| Entity | O | 0 | 0 | 0 | 0 | 0 | 0 |
| | C | 5 | 0 | 1 | 0 | 0 | 0 |
| | G | 6 | 0 | 0 | 0 | 0 | 0 |
| | M | 0 | 0 | 2 | 0 | 1 | 0 |
| Platform / device | O | 0 | 0 | 0 | 0 | 0 | 0 |
| | C | 7 | 0 | 0 | 1 | 0 | 1 |
| | G | 0 | 0 | 0 | 0 | 0 | 0 |
| | M | 0 | 0 | 0 | 0 | 0 | 1 |
| Link / path | O | 0 | 0 | 0 | 0 | 0 | 0 |
| | C | 0 | 0 | 0 | 0 | 0 | 0 |
| | G | 0 | 0 | 0 | 0 | 0 | 0 |
| | M | 0 | 0 | 1 | 0 | 0 | 0 |
| Data product | O | 0 | 4 | 0 | 2 | 0 | 0 |
| | C | 0 | 4 | 2 | 8 | 1 | 6 |
| | G | 0 | 0 | 0 | 8 | 0 | 6 |
| | M | 0 | 0 | 0 | 0 | 4 | 0 |

O: observation C: commitment/provenance G: governance M: mission use

**Evidence profiles (overlap): L0 4 L1 6 L2 13 L3 11 L4 0 L5 0**

Cells count distinct studies. Zero means no coded claim in this sample.

**Figure 3.** OAL framework and its application to 23 full-text mechanism studies. (a) The object, assurance, and lifecycle dimensions. (b) Coverage map: rows are object–assurance pairs; columns are lifecycle stages. Numbers count distinct studies per cell. The evidence-profile summary counts studies reporting each form; profiles may contain several forms. Zero denotes no coded claim in this sample, not a global research absence.

[59,60]; an attested-channel result and an embedded model-integrity result support different measured-state claims. A consumer needs the reference-value version, measurement coverage, issue time, and appraisal policy before using either result. These distinctions make it possible to combine the papers without treating the word "attestation" as one uniform operational guarantee.

Provenance studies similarly differ in what has been exercised. The sensing platform and Hadoop integration implement record capture, integrity checks, and retrieval [63,64]; the health-IoT work implements provenance alongside federated-learning applications [67]. Food-provenance experiments include task-specific forecasting and reported ledger evaluation [66]. Decision provenance and the digital-twin architecture describe broader chains of organizational and physical dependencies [56,68]. OAL preserves those broader conceptual claims while keeping their evidence categories distinct from measured implementation behavior. A prototype's successful query does not demonstrate that every upstream transformation was captured.

The map also identifies a dependency between coding and design. A study concerned with node behavior can support an entity claim about cooperation, a path claim about usable forwarding, or both when the mechanism explicitly addresses them. An endpoint's measured code state is classified under platform integrity, while the accuracy of its reported location is an observation claim about a data product. These decisions are recorded in the extraction table rather than hidden in aggregate counts. Alternative reasonable assignments can be tested by editing those records and regenerating the heatmap; they should change the interpretation of coverage before they change any claim of system effectiveness.

## 6. Evidence Dependencies, Threat Propagation, and Dynamic Reasoning

### *6.1. A Typed Object-Dependency Graph*

Let $G_t = (V_t, E_t)$ be a time-indexed directed multigraph over OAL objects and their versions or state intervals. Relations name their semantics: an entity *controls* a platform; a platform *hosts* a model or *produces* an observation; a path *transmits* a product; an entity *endorses* or *consumes* a product; and a record *commits to* another product. A transformation links its inputs, output, executing platform, responsible entity, code/model version, and event interval. In PROV, this transformation can be serialized as an activity [27].

Traversal identifies claims for reassessment when a key, endpoint, model, or input becomes disputed. A graph edge alone does not quantify causal influence or justify a probability update. Corroboration must account for shared sources, clocks, transformations, and administrative roots. Two products derived from one observation do not supply two independent observations. Conversely, a shared relay need not invalidate end-to-end signatures from otherwise independent sources.

### *6.2. Threats, Degradation, and Missing Evidence*

An analysis statement records *attacker or environmental cause; capability; affected object; propagation relation; mission consequence; control; residual uncertainty*. For example, a compromised USV can withhold UUV batches, delaying the consumer's evidence and audit; independent contacts and protected local counters may expose gaps after reconnection. Acoustic fading can create the same delivery symptom with a different cause. The response can restrict freshness-sensitive use while attribution remains unknown.

### *6.3. The CTG Reasoning Loop*

CTG organizes context, temporal state, and dependencies through evidence capture, reliability appraisal, object-state inference, dependency propagation, uncertainty expression, policy action, accountable audit, and operational feedback. Evidence capture includes

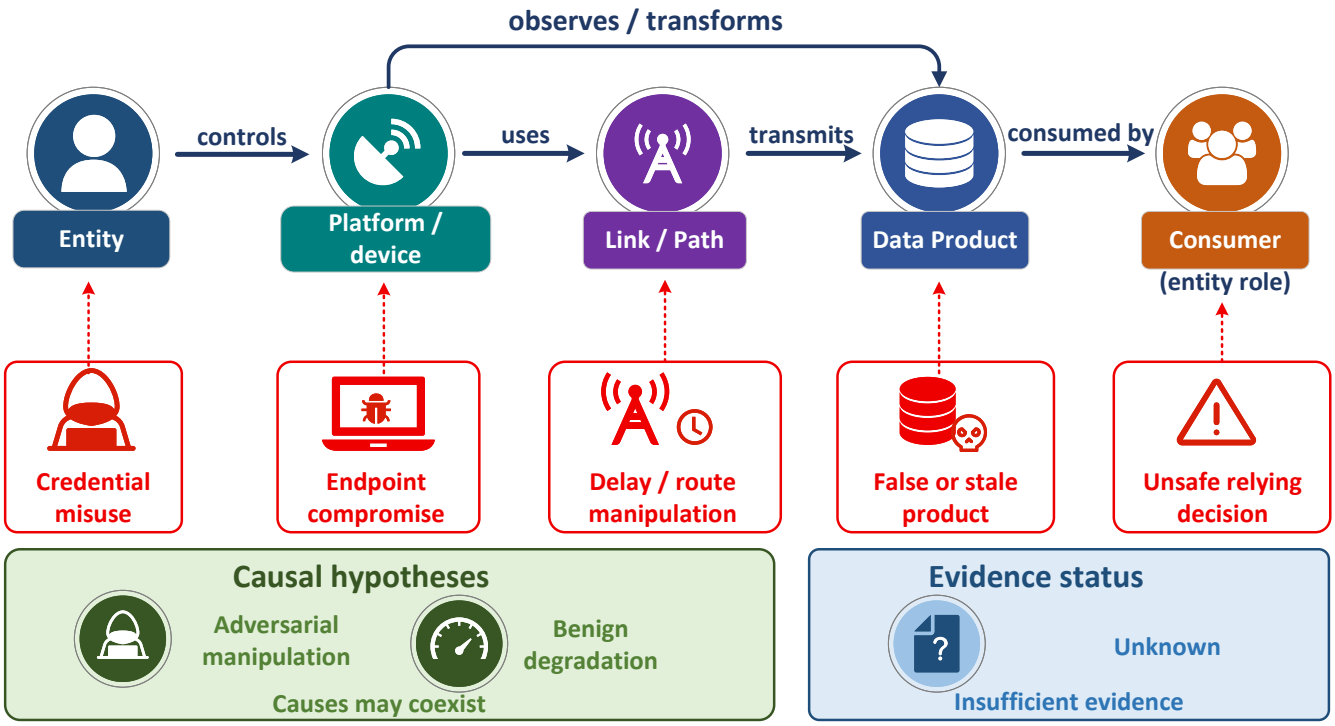



**Figure 4.** Typed dependencies and evidence states. Operational relations identify claims affected by a changed assumption. Benign degradation and adversarial intervention are causal hypotheses; unknown is an epistemic outcome.

**Table 5.** Cross-object analysis of selected failure paths.

| Initial condition | Affected dependency | Evidence and response |
|---|---|---|
| Stolen credential | Entity authority to signed product | Effective-time revocation; review signed descendants |
| Compromised endpoint | Device state to observation or transformation | Attestation scope, diagnostics, corroboration; quarantine affected outputs |
| Fading or jamming | Path availability to evidence freshness | Channel/context evidence; restrict stale use and retain uncertain cause |
| Correlated sources | Shared input to fused conclusion | Lineage and common-cause model; reduce independent-evidence count |
| Unavailable governance quorum | Shared policy to fresh global authority | Retain prior checkpoint; enforce bounded local delegation |

binding verification; propagation identifies affected descendants for reassessment under an explicit inference rule. Figure 5 depicts this information flow.

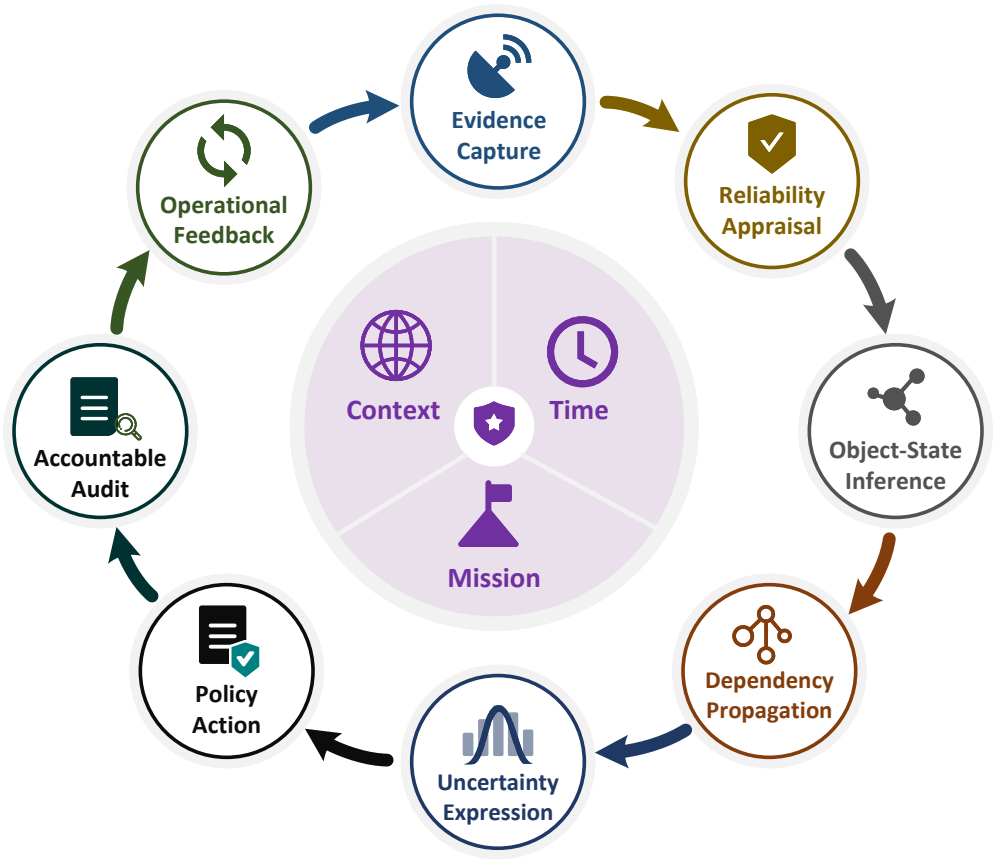



**Figure 5.** CTG appraisal and action loop, with policy, temporal state, and dependency knowledge shared across its steps.

The loop is positioned within established architectures. MAPE-K separates monitoring, analysis, planning, execution, and shared knowledge [72,73]; CTG specializes that organization for claim appraisal, event/receipt-time separation, and dependency-triggered reassessment. The JDL fusion model distinguishes object, situation, impact, and process refinement [74]; CTG adds an explicit record of which identity, provenance, governance, and temporal evidence supports reliance on those estimates. Neither relationship implies a new general-purpose control or fusion algorithm.

Implementations can use Bayesian models, subjective logic, hidden-state inference, rules, or learned classifiers. Their outputs retain the named property, context, evidence lineage, uncertainty, and permitted use. Re-estimating trust after a change requires a specified update rule, rather than an arbitrary decrement along graph edges. A hard authorization condition can coexist with a graded quality estimate; the choice to make one condition non-compensable belongs to the mission policy.

### *6.4. Partition, Revocation, and Reconciliation Trace*

The coalition in Section 3.1 has jointly approved policy version 1 and a low-consequence local delegation expiring at $t = 480$ s. At $t = 0$, the USV loses its shore connection; at $t = 240$ s, shore organizations commit a revocation of a UUV signing key, effective immediately. The UUV receives the revocation at reconnection, $t = 600$ s. Its protected clock reports event intervals, not exact event times. This scenario assumes a

trustworthy pre-compromise time/counter record for the retained events; a stolen key alone cannot establish a credible historical timestamp.

Three signed records have intervals $[114, 126]$, $[234, 246]$, and $[294, 306]$ s. On reconnection, a verifier checks their source signatures, batch inclusion, policy epoch, and the revocation history. A record whose interval ends before 240 s can be eligible for the historical audit. One beginning at or after 240 s is revoked at event time. An interval crossing 240 s is quarantined for review. Historical eligibility does not authorize a new physical action or refresh the observation.

**Table 6.** Executable replay outcomes. Ledger and policy-validating witnessed-log paths use identical evidence and authorization rules. The scalar examples use the same input features.

| Event interval (s) | OAL-organized rule / validating log | Mean score | Scalar decisions |
|---|---|---|---|
| 114–126 | Eligible historical record | 0.99 | Mean accepts; veto-aware scalar agrees with rule |
| 234–246 | Quarantine: revocation overlap | 0.89 | Mean accepts; veto-aware scalar quarantines |
| 294–306 | Revoked at event time | 0.79 | Mean accepts; veto-aware scalar rejects |

The illustrative mean averages signature validity, observation quality, lineage, cached policy match, and temporal authority, using $[1, 0.95, 1, 1, a]$ with $a = 1, 0.5, 0$ and threshold 0.75. Its compensating aggregation accepts all three records. A scalar policy with an explicit authority veto produces exactly the typed rule's decisions. Thus this example demonstrates the effect of the chosen semantics, rather than an inherent advantage of a vector representation. Comparative inference quality remains an empirical question.

During disconnection, an action at 420 s can have restricted local authority under the cached envelope even though later audit identifies a conflict with effective global revocation; at 540 s the envelope has expired and the local policy denies the action. Reconciliation retains both the locally available authority and the later global judgment. Rewriting the old decision as if the vehicle had received the revocation at 240 s would erase this distinction.

The supplied Python replay verifies Ed25519 signatures and Merkle inclusion, rejects altered records/proofs, checks all six arrival orders with duplicate delivery, and enumerates all 16 coalition signer subsets. The ledger-policy and validating-log decision functions agree throughout. These are deterministic rule checks with synthetic events; the script runs no Fabric network or consensus protocol. Section 7.3 specifies how the same application rules map to Fabric and how the constrained vehicle receives an auditable result.

## 7. Blockchain-Assisted Assurance: Selection and a Specified Pattern

### *7.1. The Decision to Use a Ledger*

Ledger adoption begins with writers, state, and authority. Wüst and Gervais organize the decision around the need to store shared state, the number and trust relationships of writers, and the availability of a trusted party [75]. Locher et al. examine which functions of a trusted third party can be replaced by distributed agreement [76]. For mission assurance, the relevant requirement is that no single organization may silently change accepted membership, policy, or revocation history, while authorized parties retain a mutually verifiable state.

In the coalition example, A alone must not replace a sensor-acceptance policy, C must not remove a previously acknowledged batch after a gateway incident, and D must not declare its own data accepted by the coalition. Four independent organizations can jointly

validate those transitions. A permissioned ledger is one implementation of this requirement. Its value is confined to such shared-governance conditions and the operational services supplied by the chosen implementation.

The strongest alternative is a transparency service whose checkpoints and state transitions are independently witnessed and checked under the same authorization rules. Such a service can satisfy the same safety requirements. A consistency-only witness detects or exposes inconsistent histories but does not by that act establish that a policy change was authorized. Adding joint policy validation closes that gap. The replay in Section 6.4 consequently finds equivalent decisions for the ledger and policy-validating log. Architecture selection then turns on implementation complexity, privacy, availability, recovery, and measured cost. No intrinsic ledger superiority is inferred.

Governance viability is also an adoption condition. Maersk and IBM discontinued TradeLens after it failed to achieve the required global industry collaboration and commercial viability, despite a technically viable platform [77]. This maritime case motivates evaluating incentives, participation, exit, and operating responsibility. It does not establish a failure of a particular consensus protocol.

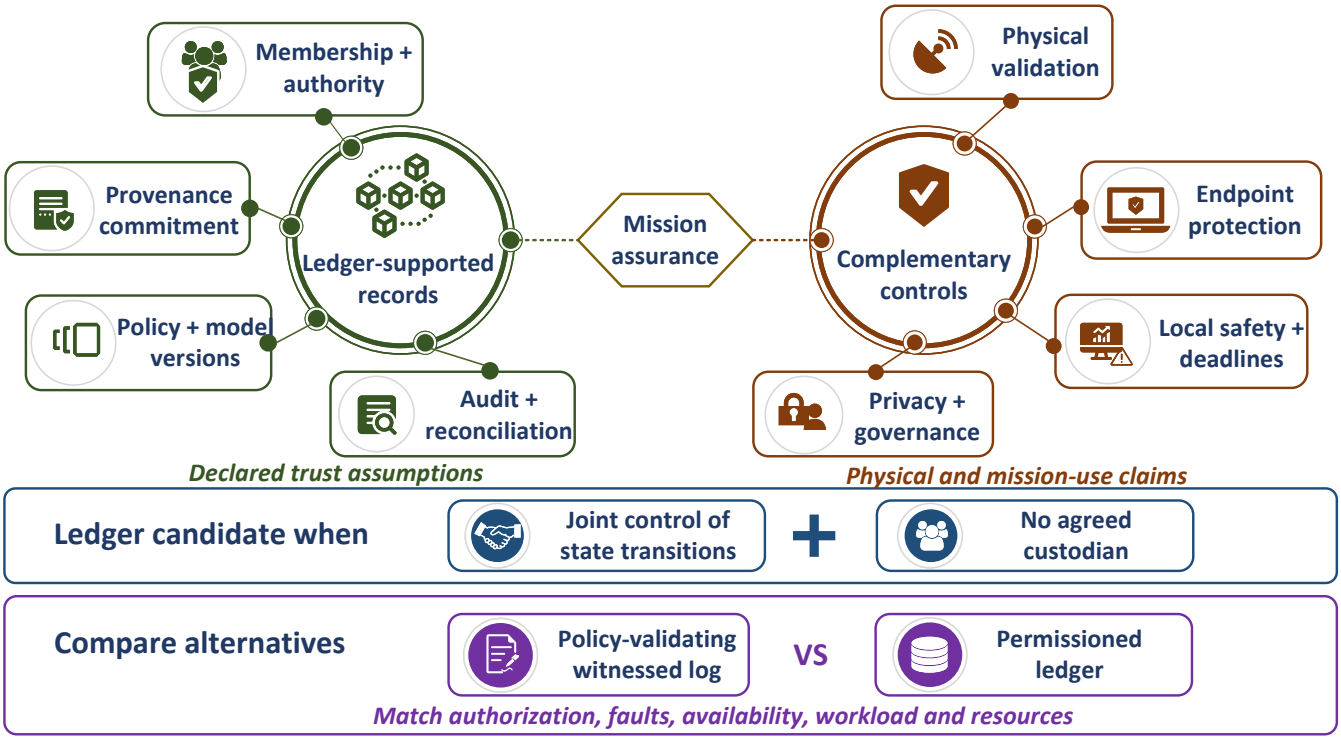



**Figure 6.** Requirement-based shared-state selection. A policy-validating witnessed log and a permissioned ledger are compared under matched governance guarantees. Physical appraisal and local safety remain application responsibilities.

### *7.2. Mission Requirements and Consensus Candidates*

A requirement-based selection process, as advocated by Švarcmajer et al. [78], specifies the fault model before choosing a protocol. Here the coalition design assumes at most one Byzantine administrative domain among four independently controlled shore nodes. Safety must hold during a USV partition; new global decisions may stop when a responsive quorum is unavailable. The illustrative reconciliation deadline is 60 s after an event becomes available for shore processing. Membership is stable during a mission and changes through a jointly approved epoch transition. Shore nodes have adequate compute and power; UUVs do not participate in consensus. Consensus traffic uses shore/radio infrastructure,

with only compact evidence commitments crossing the acoustic link. These are design requirements, rather than measurements of a deployed mission.

Table 7 compares representative families. Normal-path message complexity is a selection indicator, conditional on batching, signature aggregation, dissemination, and view changes. Fault tolerance and liveness assumptions remain decisive. In particular, a DAG organization does not remove the need to state which agreement protocol orders or certifies its data.

**Table 7.** Consensus candidates against the mission requirements. Message orders describe representative normal paths; recovery can cost more. The full requirement matrix is supplied as CSV.

| Family | Fault model and partition behavior | Latency, traffic and resource implications | Scenario fit |
|---|---|---|---|
| Raft | Crash faults; majority needed for progress | Leader replication, typically $O(n)$ exchanges per entry; low server cost; managed membership | Single accountable operator with crash-only replicas; insufficient for Byzantine coalition administrators |
| Classical BFT / SmartBFT | Byzantine faults; typically $n \geq 3f+1$ and $2f+1$ responsive votes | Multi-phase communication, commonly $O(n^2)$; timeout/view-change cost; capable stable nodes | Small independent shore coalition; selected Fabric ordering candidate |
| HotStuff | Byzantine faults; quorum and partial-synchrony liveness | Linear normal-path communication with quorum certificates and pipelining; cryptographic aggregation and leader availability matter | Larger capable coalition with measured load benefit |
| CometBFT | Byzantine faults; more than two-thirds voting power; round/locking rules | Tendermint-style propose/vote and gossip; timeouts under delay; implementation-dependent dissemination cost | Application-specific replicated state machine; not classified as HotStuff merely because both use BFT |
| Authority-based | Named validators; resilience depends on the actual protocol | Low mining cost; schedule/finality and quorum assumptions must be identified | Controlled consortium only after the concrete authority protocol meets the fault requirement |
| DAG-based (Narwhal/Tusk) | Reliable data dissemination plus an explicit BFT ordering protocol | Parallel dissemination and storage; asynchronous progress assumptions and communication cost remain | High-throughput shore service with resource evidence; no acoustic-partition shortcut |
| IoT-oriented designs | Often cluster heads, trusted managers or restricted validator sets | Reduce endpoint work by moving it to capable nodes; altered trust and availability assumptions | Gateway-mediated collection; selected only if its delegated trust matches the mission |

Raft and Byzantine protocols solve different fault problems [79–81]. CometBFT uses its own round-based consensus and gossip rules [82]; Narwhal is a dissemination/mempool component, while Tusk supplies an associated ordering protocol [83]. Authority-based UAV acquisition, as implemented in BUS, illustrates placement on capable validators [16]. Lightweight IoT proposals similarly motivate moving consensus away from constrained endpoints [84]; the optimization must be assessed together with the resulting administrative dependencies.

Fault assumptions are stated separately for infrastructure and applications. A Byzantine ordering service concerns agreement despite arbitrary behavior of replicas; it does not prevent an authorized producer from supplying a false observation. Endorsement constrains who approves a proposed state update, while the chaincode defines the deter-

ministic rule being approved. A contract can consistently implement an unsuitable rule, so policy approval also requires evidence about its mission consequence. The application must specify whether emergency authority can bypass ordinary endorsement and how that exception is recorded and reviewed. Without an explicit exception, loss of a mandatory signer is an intentional loss of fresh global authorization.

Partition tolerance likewise has two meanings that must be separated. Safety preserves a consistent accepted history when messages are delayed or unavailable. Availability concerns whether a particular request can obtain the needed ordering and application approvals before its deadline. A shore quorum may remain healthy while the USV is disconnected, allowing shore revocations to progress while a UUV continues under cached authority. If the shore quorum itself is divided, new global updates may stop even after the gateway returns. Evaluations should vary these two partition locations independently. Classical consistency/availability results motivate this distinction under network partition [85]; they do not specify the vehicle's local safety policy.

Three scenario decisions follow. A single-agency inspection mission can use an authenticated database or witnessed log, adding Raft replication for crash availability. The four-organization response mission requires Byzantine ordering plus joint application authorization, or an equivalent validating log; current Fabric SmartBFT is a concrete candidate [86]. A disconnected UUV group uses signed local records, store-and-forward transport, and bounded delegation, with no requirement for acoustic consensus. Increasing validator count only helps when it increases independent fault domains while preserving a reachable quorum. Actual latency and energy must be measured for the selected deployment.

### *7.3. A Coalition Provenance-and-Revocation Pattern*

Five recurring ledger roles are membership/authorization, evidence commitments, policy/model version coordination, accountable audit, and post-partition reconciliation. The specified pattern combines commitments and reconciliation; Table 8 places the other roles around it. The application logic is specified in the supplementary pseudocode, and its temporal and authorization rules are exercised by the replay.

**Table 8.** Ledger roles and their corresponding application requirements.

| Role | Shared state | Required complementary evidence |
|---|---|---|
| Membership | Role, delegation, key epoch | Identity proof and protected key custody |
| Commitments | Batch root, sequence, retained-object reference | Signed source records and retrievable payload |
| Version coordination | Policy/model hash and approval epoch | Model validation and deployment evidence |
| Audit | Decision, evidence links and effective-time history | Complete capture and independent review |
| Reconciliation | Received batch, gaps, conflicts, judgments | Protected clock/counter and offline-authority history |

**Placement and endorsement.** Each organization A–D hosts an endorsing peer and a separate ordering node on capable shore infrastructure. The illustrative BFT service has four orderers, tolerating one Byzantine node under its protocol assumptions. Administrative separation, rather than merely four processes, supplies the independence. The application endorsement rule is `AND(A, B, OR(C, D))`; channel configuration and chaincode-definition approval require the same coalition authority. A and B are mandatory, so application progress can stop even while three orderers remain responsive. The USV forwards and caches signed evidence, with no consensus vote. A partitioned gateway

buffers batches and retries idempotently after reconnection. UUV safety control uses its local envelope throughout.

**State and encoding.** The state keys are mission/policy/epoch, producer/key/history, delegation/identifier, batch/producer/sequence, and judgment/record/policy. A compact source record has a 154-byte fixed-width body: version (1), kind (1), mission ID (16), producer ID (16), sequence (4), event lower and upper bounds (8 each), authority epoch (4), policy hash (32), payload hash (32), and parent/root commitment (32). Integers use network byte order; time is signed seconds in the declared mission epoch. A 64-byte Ed25519 signature gives a 218-byte record. Domain-separated signing and hashing bind record type and version. A batch uses the same body layout with a Merkle root and previous-batch root, plus a two-byte record count: 156 bytes before signature and 220 bytes signed. Ten signed records are retained off chain. The released canonical JSON example submitted by a gateway is exactly 525 application bytes, excluding the file's trailing newline.

**Deterministic chaincode operations.** `CommitBatch` reads the mission epoch, gateway role, and last batch sequence; verifies the compact gateway signature and policy binding; rejects an inconsistent sequence or duplicate identifier with different content; and records the root. An identical retry returns the existing result. This commits an auditable batch, not a claim that every leaf is fit for use. `JudgeRecord` verifies a supplied source signature and inclusion proof, reads the producer's historical key status and delegation, and writes a separate eligible/revoked/quarantined judgment using the interval rule in Section 6.4. `UpdatePolicy` and `RevokeKey` require the coalition endorsement rule and monotone epochs; revocation records preserve their effective time. All state read for a decision enters the Fabric read set, allowing validation to reject stale concurrent updates. Chaincode obtains no external clock or oracle value during execution; authorized submitted evidence supplies those inputs.

**Commit status and constrained verification.** Fabric ordering precedes peer validation, so block inclusion is not proof of a valid world-state update [62,87]. After observing a valid commit, three distinct organizations satisfying A+B+(C or D) sign an application receipt containing the mission, policy/authority epoch, batch root, sequence, event span, and count. Receipt-signing code must check the commit validation status and the expected application state. A UUV checks authorized signer identities against its pinned epoch, signatures, monotone checkpoint sequence, source-record signature, and the Merkle branch. For ten leaves padded to 16, the branch has four 32-byte hashes. These checks establish a jointly attested application checkpoint and inclusion under pinned keys; they do not constitute native verification of all Fabric consensus and peer-validation execution. A new epoch requires an authenticated transition from the pinned epoch. The receipt is an additional application trust boundary that must be implemented and tested.

### *7.4. Communication Budget and Partition Sensitivity*

The receipt and retrieval paths serve different evidence needs. A vehicle receiving a checkpoint can verify that independent organizations acknowledged the committed root under a pinned authority epoch. To appraise an observation, it still obtains the relevant leaf, its inclusion path, and any required source or transformation evidence. A commitment-only contact therefore supports a narrower claim than complete evidence delivery. Retrieval policy can prioritize recent high-consequence products and defer bulk imagery until a higher-bandwidth contact, while recording which required inputs remain unavailable. A failed retrieval is represented as missing evidence rather than a successful integrity check.

Clock and key transitions also affect the usable lifetime of a proof. A cached checkpoint remains evidence about a historical commitment after a new policy is issued, but its author-

ity to support a new decision depends on the relying policy and authenticated transition history. A vehicle with an uncertain clock cannot establish that an expiry condition still holds by inspecting a ledger root. It needs a bounded time interval or a more restrictive fallback. Similarly, a fresh receipt cannot make an untrustworthy acquisition timestamp credible. These requirements bind the application-level proof to the same temporal assumptions used in the replay and define what a constrained verifier must do when those assumptions fail.

Ed25519 signatures contain 64 bytes [88]; ML-DSA-44 signatures contain 2420 bytes [89]. These are alternative compact application profiles, not claims that either encoding is the default Fabric transaction-signature format. Fabric enrollment certificates, endorsements, proposal responses, read/write sets, and envelopes add deployment-dependent bytes. Only the gateway transports those full transactions on the shore network. Their total size and resource cost require a configured implementation; the 525-byte JSON value is not presented as a measured Fabric transaction size.

For a message of $B$ bytes and net payload rate $R$ bit/s, the serialization lower bound is $8B/R$. Table 9 uses 100, 1000, and 5000 bit/s as scenario inputs. It excludes propagation, contact waiting, modem framing, BPSec overhead, loss, retries, and certificate/epoch updates. Those factors increase end-to-end delay and must be added for a specified modem and transport configuration.

**Table 9.** Application-message serialization lower bounds. All sizes are computed from the released encoding; ML-DSA values substitute the standardized signature length.

| Message | Bytes | 100 bit/s | 1 kbit/s | 5 kbit/s |
|---|---|---|---|---|
| Ed25519 source record | 218 | 17.44 s | 1.744 s | 0.349 s |
| Ed25519 batch commitment (10 records) | 220 | 17.60 s | 1.760 s | 0.352 s |
| Ten complete Ed25519 records | 2180 | 174.40 s | 17.440 s | 3.488 s |
| ML-DSA-44 batch commitment | 2576 | 206.08 s | 20.608 s | 4.122 s |
| Ed25519 record + branch + joint receipt | 697 | 55.76 s | 5.576 s | 1.115 s |

The 697-byte downlink comprises a 218-byte source record, 128-byte branch, two-byte leaf index, 156-byte checkpoint, one-byte signer bitmap, and three 64-byte signatures. At one ten-record batch per minute, the 220-byte commitment needs 29.3 net bit/s on average, whereas sending all ten signed records needs 290.7 bit/s. Commitment-only delivery therefore fits some contacts that cannot carry full evidence, at the cost of later retrieval and incomplete immediate appraisal. Sending one ten-record ML-DSA-44 batch commitment at the same cadence requires 343.5 bit/s. Public keys are assumed provisioned; key distribution and post-quantum verification time remain separate costs.

The replay also enumerates 63 combinations of validator count (4, 7, 10), responsive count (quorum minus one, quorum, all), and partition duration (0, 10, 30, 58, 59, 60, 120 s). With assumed post-contact delays of 2 s for a ledger path and 1 s for a validating-log path, a 60 s deadline permits partitions of at most 58 and 59 s, respectively, when a quorum is reachable. Below quorum, neither completes a fresh joint decision. This sensitivity calculation exhibits a regime in which the log is faster and no ledger safety gain exists; its delays are illustrative inputs, not benchmarks. Application endorsement availability must be checked in addition to ordering quorum. The result motivates measuring both architectures under the same workload, governance, and fault domains.

## 8. Evaluation Practices and Evidence Maturity

### *8.1. A Claim-Specific Evidence Profile*

L0–L5 record how a claim is supported: concept, analysis, simulation/data, prototype, field trial, and sustained operation. Figure 7 presents these complementary evidence forms using schematic representations. Formal soundness, experimental control, environmental representativeness, and operational duration remain separate assessment dimensions.

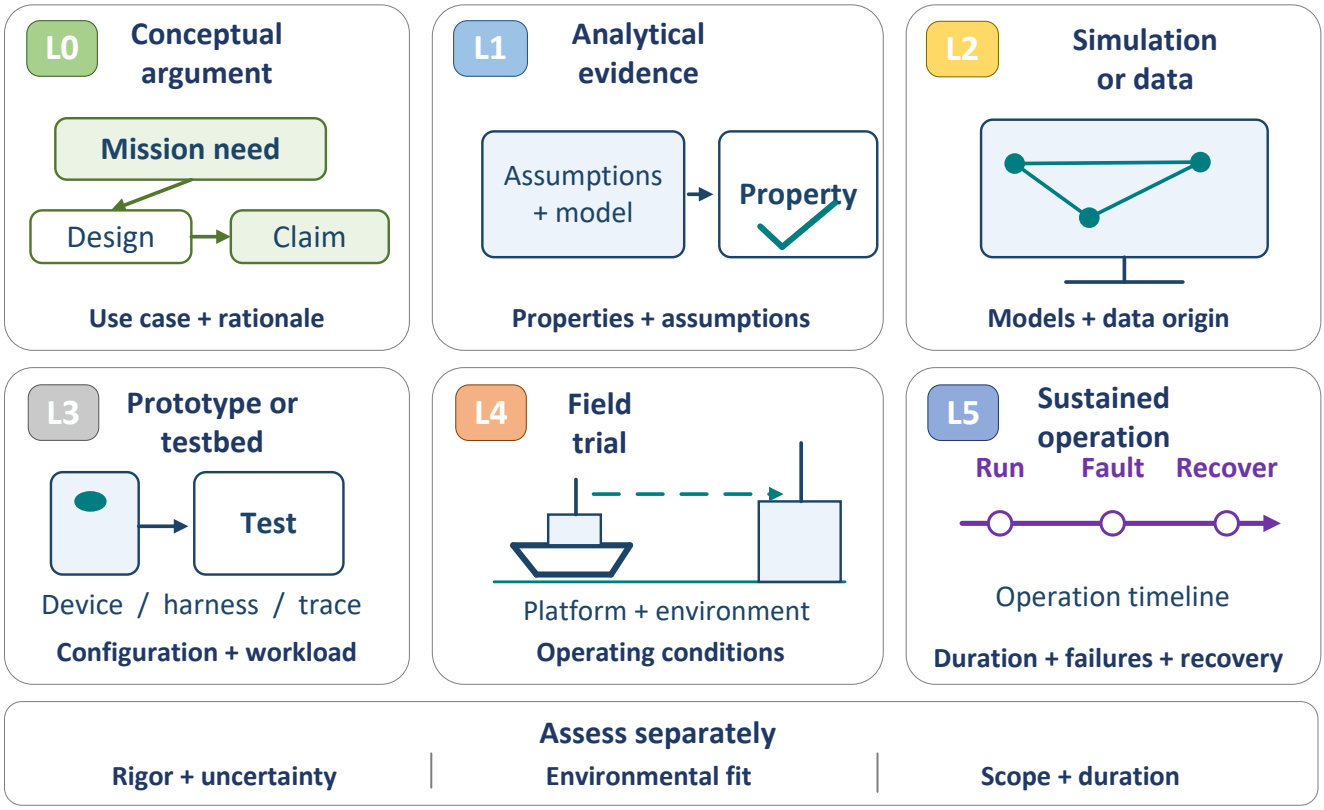



**Figure 7.** Complementary evidence forms and the scope each can support. The schematics are drawn as native Visio shapes; no photographic or generated-image assets are used in this figure.

Within the 23-study sample, four profiles include L0, six L1, thirteen L2, and eleven L3; none is coded L4 or L5 for the extracted claims. Profiles overlap. BUS includes UAV and IoT hardware, but its experimental description excludes physical channel parameters and does not establish an integrated maritime field setting. GNSS validation on real-world recordings supplies L2 evidence for the evaluated detection task. These distinctions prevent component implementations and recorded field data from being treated as demonstrations of whole-mission assurance.

### *8.2. Datasets, Simulators, and Testbeds*

Public resources support complementary evaluation tasks. CICIoT2023, TON_IoT, and Edge-IIoTset supply attack and telemetry data for classification and robustness tests [90–92]. SeaDronesSee supplies maritime visual data [93]. AirSim, UUV Simulator, and Aqua-Sim NG address vehicle/sensor simulation, underwater dynamics, and underwater communication, respectively [94–96]. Their use in one experiment requires compatible clocks, geometry, mobility, and event interfaces, together with a separately validated mission-loss model.

**Table 10.** Evaluation resources and required transfer checks.

| Resource | Useful evidence | Required complement |
|---|---|---|
| IoT attack/telemetry datasets | Detection, calibration, missingness and false isolation | Maritime topology, benign degradation and mission consequences |
| SeaDronesSee / GNSS recordings | Visual detection or signal appraisal | Independent ground truth for the stated physical claim |
| AirSim / UUV Simulator | Dynamics, sensors, disturbance and control | Field validity of the configured models |
| Aqua-Sim NG | Path delay, loss, energy and contact behavior | Physical observation and governance modules |
| Ledger and witnessed-log testbeds | Validation, bytes, latency, recovery and membership changes | Matched guarantees, evidence workload and fault domains |
| Multi-organization field exercise | Integrated decisions, recovery and audit | Repeated scenarios and explicit operational duration |

### *8.3. Metrics, Baselines, and Reproducibility*

An evaluation should measure inference error and uncertainty, mission consequences, network/resource cost, shared-state behavior, and accountability. Useful outcomes include unsafe acceptance, false isolation of a critical relay, task completion, safety-constraint violations, deadline misses, bytes per evidence event, energy, endorsement/commit latency, time to identify affected descendants, and the fraction of decisions with retrievable provenance. Reporting only throughput omits the consequence of a timely but unsuitable decision.

Scenario factors include compromised entities versus devices, independent versus correlated observations, fading versus jamming, fresh versus stale attestation, lost contacts, collusion, and different authority-expiry periods. Ground truth distinguishes benign degradation, adversarial intervention, and their coexistence; unknown is an appraisal outcome. Training, tuning, calibration, and evaluation use separate data, with splits respecting shared missions, platforms, and recording episodes. Paired scenarios and uncertainty intervals over independent runs support comparative claims. Undelivered events and deadline misses remain in the latency/availability analysis.

Baselines receive the same raw evidence, action set, mission objectives, and tuning budget. Appraisal comparisons include a scalar model with explicit authorization vetoes, and component ablations for lineage, shared causes, and uncertainty. Storage/governance comparisons include an authenticated database, an independently witnessed consistency log, and a policy-validating jointly cosigned log. The strongest baseline matches administrative independence, quorum, retention, and effective-time semantics. Differences in guarantees are reported before interpreting differences in cost.

A reproducible report supplies topology, fault domains, channel/contact models, software versions, payload encodings, cryptographic suites, endorsement rules, policy/model versions, workload, seeds, ground truth, and mission-loss weights. The present supplement provides the claim records, heatmap aggregation, replay inputs/results, byte calculations, and chaincode specification. It enables reproduction of the analytical illustration; deploying and benchmarking the specified Fabric and witnessed-log implementations is a subsequent evaluation stage.

Evaluation should also distinguish collection completeness from record consistency. A test can maintain an independent event oracle, then inject omission before signing, deletion of an unwitnessed suffix, relay withholding after commitment, and loss of an off-chain

object after its root has been accepted. These failures have different observability. A retained receipt exposes the absence of an acknowledged batch; it cannot reveal an event never captured by any trusted observer. Metrics should therefore report the denominator of events expected to be captured, those acknowledged by a witness, those committed under joint policy, and those whose payloads remain retrievable. This separates preservation, detection, and recovery rather than treating all three as one audit-success rate.

Recovery workloads need similar precision. A returned gateway may carry valid old observations, revoked-key submissions, conflicting sequence numbers, and records generated under different policy versions. Benchmarking only clean insertion misses this mix. Experiments should vary backlog age and size, compromised-key cutoff, duplicate rate, epoch transitions, proof retrieval, and the proportion requiring human review. Report both automated processing time and the time until the application can act on a resolved judgment. Shared-state throughput can remain high while semantic reconciliation is blocked by missing evidence or unavailable mandatory endorsers.

Finally, comparison of inference policies should preserve the cost of their actions. Quarantining every uncertain observation can reduce unsafe acceptance while preventing useful missions. A declared loss model assigns costs to unsafe reliance, unnecessary isolation, deferred tasks, corroboration requests, and human review. Results should show sensitivity to those weights and to the prevalence of attacks and benign degradation. Calibration and decision loss answer different questions: a well-calibrated estimate may still be used by an unsuitable policy, while a conservative rule can make safe decisions without producing probabilities. Keeping the two evaluations separate follows the appraisal/decision interface in Equation (1).

## 9. Discussion and Research Agenda

### *9.1. Integration, Governance, and Privacy*

OAL makes the composition problem explicit: provenance concerns history, observation validity concerns the world, shared-state consistency concerns agreement, and mission suitability concerns consequence. A policy specifies how these claims combine for a given use. Dependency analysis identifies where a changed key, model, input, or reference value invalidates an assumption. It also exposes architectural common causes: adding several UUVs does not remove their dependence on one gateway, clock, or operator.

Consortium design must specify admission, contract/model approval, emergency authority, disputes, member exit, and audit responsibility. Maritime adoption research and the TradeLens experience show why participation and incentives belong in this analysis [20,77,97]. The endorsement rule in the example preserves joint control while creating a measurable availability cost through its mandatory signers. Additional replicas under one owner do not create additional independent governance domains.

Cross-sector adoption studies also distinguish the claimed benefits of traceability from the organizational changes needed to realize them [98,99]. For a coalition, the unit of accountability is the participating authority and its commitments, not simply the number of ledger nodes. A useful deployment agreement assigns evidence retention, disclosure permissions, service maintenance, and incident-response duties before a mission. It defines what remains available when an organization leaves and which signatures continue to support historical judgments after key rotation. These provisions can be implemented with a validating log or ledger; they form part of the matched governance assumptions in a fair comparison.

Privacy depends on more than payload encryption. Stable identifiers, route custody, endorsement timing, and model-use records can reveal vehicle activity and organizational relationships [100,101]. Data minimization, selective disclosure, retention, and controlled

identity resolution therefore accompany auditability. A revocation invalidates future reliance under a policy; it does not erase a retained commitment. Human override is a separately authorized, recorded transition with its evidence, rationale, and consequence.

### *9.2. Falsifiable Research Propositions*

The following propositions define future comparative tests. Each experiment prespecifies its mission distribution, minimum useful effect, primary outcome, and analysis rule. Null and adverse results identify conditions where the additional structure provides no benefit.

**P1 (claim organization).** Test whether a separately represented assurance vector reduces held-out unsafe acceptance by a prespecified amount at matched false-isolation and latency budgets, relative to a tuned scalar policy receiving the same complete evidence. The scalar comparator includes non-compensable authorization/freshness gates and may encode interactions. Both use the same actions and equal tuning budgets. The proposition fails if the matched comparator performs equally well or better; the illustrative veto-aware scalar in Table 6 is one such equality case. Representation alone does not guarantee improvement.

**P2 (dependency modeling).** Under prespecified shared-source error and collusion rates, test whether a model with common-cause variables improves a proper scoring rule and mission loss relative to an equally tuned independence model. Include correct, incomplete, and erroneous lineage conditions. The proposition fails when model misspecification or acquisition cost offsets the gain, or when observations are sufficiently independent that no useful improvement remains.

**P3 (diagnosis and unknown status).** Compare a model retaining benign, adversarial, and unresolved hypotheses with a binary appraisal model under mixed degradation, attacks, and evidence outages. Both can choose the same actions, including abstention. Separate diagnosis and abstention ablations, and measure loss from unsafe acceptance, false isolation, corroboration delay, and lost availability. The proposition fails if the added diagnostic state increases loss or yields no useful improvement.

**P4 (offline authority).** Vary delegation scope, expiry, clock uncertainty, partition duration, and revocation time. Test whether bounded delegation offers a useful trade-off between tasks completed before deadline and actions conflicting with effective global policy, compared with fresh-global-only authorization and unrestricted offline autonomy. Report locally authorized actions and retrospectively conflicting actions separately. Failure includes operating regimes where the bounded policy is dominated on both outcomes.

**P5 (conditional ledger value).** Compare an implemented permissioned ledger with a policy-validating independently witnessed log under equal membership, fault domains, validation rules, retention, and workload. Vary $n$, responsive quorum, partition duration $D$, event rate, and recovery backlog. For deadline $\tau$, estimate the boundary $D + L(n, \lambda) \leq \tau$, where $L$ is measured post-contact completion delay at load $\lambda$. Test for a prespecified improvement in timely, mutually verifiable decisions or recovery completeness at a fixed byte/energy budget. Equal safety with higher ledger cost, lower availability from mandatory endorsers, or no improvement over the validating log refutes the proposed advantage in that regime. The 58/59 s boundary in Section 7.4 illustrates this test with assumed delays; it is not its empirical result.

**P6 (evaluation coverage).** Select configurations from identical candidate sets and equal evaluation budgets using either transaction/network metrics alone or those metrics plus observation, endpoint, governance, and mission outcomes. Compare held-out mission loss under fixed selection rules. The proposition fails if the broader evaluation does not improve selection after its additional acquisition and interpretation costs are included.

### *9.3. Assurance Cases and Maritime Requirements*

OAL supplies the content of an assurance argument; existing notations supply its structure. GSN represents claims, strategies, context, assumptions, and evidence, while SACM provides a structured assurance-case metamodel [102,103]. UL 4600 supplies safety-case requirements for autonomous products rather than a diagram notation [104]. OAL can annotate their claims with object/version, assurance layer, lifecycle stage, event interval, and evidence profile. It does not replace those standards or establish compliance with them.

For a claim that a fused product is suitable for route replanning, subclaims identify credible observations, protected derivation, valid authority, current policy, and acceptable uncertainty. A failed dependency is a defeater or evidence gap requiring review. The argument retains observation, issue, receipt, and expiry times, and identifies decisions made under cached authority. A corrected input or revoked key triggers review of affected descendants while preserving independently supported products. Validation should measure critical evidence gaps, seeded invalid arguments, reviewer error, and time to identify affected products.

Maritime requirements add a specific application boundary. IMO's interim guidelines for MASS trials and regulatory scoping exercise provide trial and regulatory context [105,106]. As of the revision date, the non-mandatory MASS Code had been adopted in May 2026 and taken effect on 1 July 2026 [107]. DNV-CG-0264 addresses autonomous and remotely operated ships [108]. An assurance case must identify the vessel and operation to which these instruments apply, the responsible authority, operational design conditions, human/remote-control responsibilities, and evidence of safe fallback. Their maritime scope cannot be transferred wholesale to UAV or underwater operations. A ledger record can trace an approval or operating limit; the substantive safety evidence and applicable approval process remain separate.

### *9.4. Toward a Cross-Domain Unmanned Trust Profile*

The proposed Cross-Domain Unmanned Trust Profile (CDUTP) is an interoperability agenda for versioned claims and evidence references. It can reuse RATS roles, W3C PROV relations, zero-trust policy concepts, and secure-update manifests [24,26,28,109]. UAS and maritime service specifications supply communications context [110,111]. A profile instance records object/type, scoped claim, issuer and relying role, evidence reference, event and receipt times, expiry, uncertainty, permitted use, policy/model/schema versions, disclosure conditions, dependencies, and revocation binding.

The compact encoding in Section 7.3 is one illustrative transport profile. Interoperability requires version-specific semantic mappings and tests, including bounded clocks, key rotation, missing evidence, and disconnected issuance. Independent implementations can then be compared in plugfests and multi-organization exercises. The key research question is whether common semantics reduce decision and audit errors at an acceptable communication and governance cost.

## 10. Limitations

The narrative selection and 23-study analytical sample provide an inspectable synthesis rather than exhaustive coverage. The original discovery procedure retained candidate records but not all historical rejection decisions; the revision has not reconstructed missing screening data. Fourteen original mechanism candidates remain outside full-text numerical coding, and a single extraction pass provides no independent inter-rater estimate. The heatmap is sensitive to inclusion, object/layer interpretation, and claim granularity. Empty cells and the absence of L4/L5 codes describe this sample only. Future synthesis should

expand retrieval, independently duplicate coding, and report disagreement and sensitivity to alternative assignments.

OAL and CTG organize claims, dependencies, and policy interfaces. They do not supply a general inference theorem, calibrated trust model, or proof that every consequence-relevant dependency has been identified. Correlation estimates, causal interpretations, model transfer, and evidence sufficiency require application-specific validation. A scalar policy can preserve the same constraints as a structured representation, as the worked case demonstrates. Claims of improved mission outcomes therefore require the matched comparisons in Section 9.2.

The revision case uses synthetic events, a fixed coalition, protected-clock assumptions, and declared revocation semantics. It checks signatures, inclusion, deterministic policy rules, and byte lengths. It does not execute Fabric, measure consensus or acoustic performance, implement a production receipt service, or validate integrated UAV–USV–UUV safety. Full Fabric transaction sizes, provisioning overhead, cryptographic energy, clock rollback resistance, modem behavior, receipt issuance, and operator response remain implementation measurements. Communication times are serialization lower bounds; quorum sensitivities use assumed delays. These distinctions define what the reproducible illustration establishes.

Cryptographic evidence depends on protected keys, credible capture, retrievable payloads, and authenticated time or counters. An attacker who controls an unprotected key and clock can fabricate apparently historical records. A Merkle root does not recover withheld data or prove that an event was captured. Joint validation depends on genuinely independent administrative domains and an available application endorsement set. Governance deadlock, collusion beyond the fault model, and disputed physical facts require procedures outside consensus.

Transfer to maritime operations requires environmental and regulatory assessment for the actual vehicle, mission, and jurisdiction. Public datasets and component prototypes do not determine the consequences of rare faults in a deployed coalition. Field studies should evaluate privacy, retained metadata, human authority, recovery, and sustained organizational participation alongside technical behavior.

## 11. Conclusions

Mission trust connects evidence about physical observations, protected records, jointly governed state, and fitness for a particular use. OAL organizes these claims by object and lifecycle, while a typed dependency graph and CTG loop identify when changes require reassessment. Applying the framework to 23 full-text studies makes the sampled concentrations and incomplete transitions inspectable through claim records and a coverage heatmap.

The partition–revocation case specifies coalition roles, a Fabric-oriented application pattern, constrained verification, and a reproducible communication budget. Its executable replay preserves the distinction between local authority, event-time revocation, and later audit. A policy-validating witnessed log produces the same decisions under matched rules. Ledger value is therefore conditional on shared-governance requirements and the measured cost of implementing them.

The resulting research agenda calls for independently reproducible coding, matched architecture and appraisal baselines, realistic partition and recovery tests, and integrated mission evidence. These steps can determine where the proposed organization improves decisions and accountability, and where a simpler implementation satisfies the same claims.

**Supplementary Materials:** Supplementary materials contain the analytical CSV files, codebook, protocol specification, and Python reproduction scripts described in Sections 2.2, 6.4, and 7.4.

**Author Contributions:** Author contribution details are withheld for double-anonymous review and will be supplied after peer review using the CRediT taxonomy. All authors will confirm that they have read and agreed to the published version of the manuscript.

**Funding:** Funding information is withheld for double-anonymous review and will be supplied after peer review.

**Institutional Review Board Statement:** Not applicable.

**Informed Consent Statement:** Not applicable.

**Data Availability Statement:** The supplementary materials provide the OAL claim mapping, original-citation dispositions, query list, consensus-requirement matrix, pattern specification, executable replay, computed results, editable Visio figures, and PDF export script. Source publications are cited in the references; copyrighted full texts are not redistributed.

**Acknowledgments:** The authors thank those who provided constructive feedback during the preparation of this manuscript.

**Conflicts of Interest:** The authors declare no conflicts of interest.

## Abbreviations

| | |
|---|---|
| AUV | Autonomous underwater vehicle |
| CDUTP | Cross-Domain Unmanned Trust Profile |
| CTG | Context–temporal–graph |
| DLT | Distributed ledger technology |
| OAL | Object–Assurance–Lifecycle |
| RATS | Remote ATtestation procedureS |
| UAS | Unmanned aerial system |
| UAV | Unmanned aerial vehicle |
| USV | Unmanned surface vehicle |
| UUV | Unmanned underwater vehicle |
| ZTA | Zero trust architecture |

## References

1. Li, J.; Zhang, G.; Jiang, C.; Zhang, W. A survey of maritime unmanned search system: Theory, applications and future directions. *Ocean Engineering* **2023**, *285*, 115359. https://doi.org/10.1016/j.oceaneng.2023.115359.
2. Wei, W.; Wang, J.; Fang, Z.; Chen, J.; Ren, Y.; Dong, Y. 3U: Joint Design of UAV-USV-UUV Networks for Cooperative Target Hunting. *IEEE Transactions on Vehicular Technology* **2023**, *72*, 4085–4090. https://doi.org/10.1109/tvt.2022.3220856.
3. Pang, W.; Zhu, D.; Chen, M.; Wang, Z.; Wang, B. A cooperative heterogeneous formation control and motion planning design of 3U: UAV-USV-UUV systems. *Ocean Engineering* **2026**, *347*, 123867. https://doi.org/10.1016/j.oceaneng.2025.123867.
4. Hayat, S.; Yanmaz, E.; Muzaffar, R. Survey on Unmanned Aerial Vehicle Networks for Civil Applications: A Communications Viewpoint. *IEEE Communications Surveys & Tutorials* **2016**, *18*, 2624–2661. https://doi.org/10.1109/comst.2016.2560343.
5. Shakhatreh, H.; Sawalmeh, A.H.; Al-Fuqaha, A.; Dou, Z.; Almaita, E.; Khalil, I.; Othman, N.S.; Khreishah, A.; Guizani, M. Unmanned Aerial Vehicles (UAVs): A Survey on Civil Applications and Key Research Challenges. *IEEE Access* **2019**, *7*, 48572–48634. https://doi.org/10.1109/access.2019.2909530.
6. Bae, I.; Hong, J. Survey on the Developments of Unmanned Marine Vehicles: Intelligence and Cooperation. *Sensors* **2023**, *23*, 4643. https://doi.org/10.3390/s23104643.
7. González-García, J.; Gómez-Espinosa, A.; Cuan-Urquizo, E.; García-Valdovinos, L.G.; Salgado-Jiménez, T.; Cabello, J.A.E. Autonomous Underwater Vehicles: Localization, Navigation, and Communication for Collaborative Missions. *Applied Sciences* **2020**, *10*, 1256. https://doi.org/10.3390/app10041256.
8. Mekdad, Y.; Aris, A.; Babun, L.; Fergougui, A.E.; Conti, M.; Lazzeretti, R.; Uluagac, A.S. A survey on security and privacy issues of UAVs. *Computer Networks* **2023**, *224*, 109626. https://doi.org/10.1016/j.comnet.2023.109626.
9. Wang, X.; Zhao, Z.; Yi, L.; Ning, Z.; Guo, L.; Yu, F.R.; Guo, S. A Survey on Security of UAV Swarm Networks: Attacks and Countermeasures. *ACM Computing Surveys* **2025**, *57*, 1–37. https://doi.org/10.1145/3703625.

10. Hadi, H.J.; Cao, Y.; Nisa, K.U.; Jamil, A.M.; Ni, Q. A comprehensive survey on security, privacy issues and emerging defence technologies for UAVs. *Journal of Network and Computer Applications* **2023**, *213*, 103607. https://doi.org/10.1016/j.jnca.2023.103607.
11. Jiang, S. On Securing Underwater Acoustic Networks: A Survey. *IEEE Communications Surveys & Tutorials* **2019**, *21*, 729–752. https://doi.org/10.1109/comst.2018.2864127.
12. Yisa, A.G.; Dargahi, T.; Belguith, S.; Hammoudeh, M. Security challenges of Internet of Underwater Things: A systematic literature review. *Transactions on Emerging Telecommunications Technologies* **2021**, *32*. https://doi.org/10.1002/ett.4203.
13. Han, P.; Sui, A.; Wu, J. Identity Management and Authentication of a UAV Swarm Based on a Blockchain. *Applied Sciences* **2022**, *12*, 10524. https://doi.org/10.3390/app122010524.
14. Tan, Y.; Wang, J.; Liu, J.; Kato, N. Blockchain-Assisted Distributed and Lightweight Authentication Service for Industrial Unmanned Aerial Vehicles. *IEEE Internet of Things Journal* **2022**, *9*, 16928–16940. https://doi.org/10.1109/jiot.2022.3142251.
15. Karmakar, R.; Kaddoum, G.; Akhrif, O. A Blockchain-Based Distributed and Intelligent Clustering-Enabled Authentication Protocol for UAV Swarms. *IEEE Transactions on Mobile Computing* **2024**, *23*, 6178–6195. https://doi.org/10.1109/tmc.2023.3319544.
16. Islam, A.; Shin, S.Y. BUS: A Blockchain-Enabled Data Acquisition Scheme With the Assistance of UAV Swarm in Internet of Things. *IEEE Access* **2019**, *7*, 103231–103249. https://doi.org/10.1109/access.2019.2930774.
17. Hafeez, S.; Khan, A.R.; Al-Quraan, M.M.; Mohjazi, L.; Zoha, A.; Imran, M.A.; Sun, Y. Blockchain-Assisted UAV Communication Systems: A Comprehensive Survey. *IEEE Open Journal of Vehicular Technology* **2023**, *4*, 558–580. https://doi.org/10.1109/ojvt.2023.3295208.
18. Alladi, T.; Chamola, V.; Sahu, N.; Guizani, M. Applications of blockchain in unmanned aerial vehicles: A review. *Vehicular Communications* **2020**, *23*, 100249. https://doi.org/10.1016/j.vehcom.2020.100249.
19. Mehta, P.; Gupta, R.; Tanwar, S. Blockchain envisioned UAV networks: Challenges, solutions, and comparisons. *Computer Communications* **2020**, *151*, 518–538. https://doi.org/10.1016/j.comcom.2020.01.023.
20. Farah, M.B.; Ahmed, Y.; Mahmoud, H.; Shah, S.A.; Al-kadri, M.O.; Taramonli, S.; Bellekens, X.; Abozariba, R.; Idrissi, M.; Aneiba, A. A survey on blockchain technology in the maritime industry: Challenges and future perspectives. *Future Generation Computer Systems* **2024**, *157*, 618–637. https://doi.org/10.1016/j.future.2024.03.046.
21. Li, L.; Zhou, H. A survey of blockchain with applications in maritime and shipping industry. *Information Systems and e-Business Management* **2021**, *19*, 789–807. https://doi.org/10.1007/s10257-020-00480-6.
22. Caldarelli, G. Understanding the Blockchain Oracle Problem: A Call for Action. *Information* **2020**, *11*, 509. https://doi.org/10.3390/info11110509.
23. Al-Breiki, H.; Rehman, M.H.U.; Salah, K.; Svetinovic, D. Trustworthy Blockchain Oracles: Review, Comparison, and Open Research Challenges. *IEEE Access* **2020**, *8*, 85675–85685. https://doi.org/10.1109/ACCESS.2020.2992698.
24. Rose, S.; Borchert, O.; Mitchell, S.; Connelly, S. Zero Trust Architecture. Technical report, National Institute of Standards and Technology, 2020. https://doi.org/10.6028/nist.sp.800-207.
25. Syed, N.F.; Shah, S.W.; Shaghaghi, A.; Anwar, A.; Baig, Z.; Doss, R. Zero Trust Architecture (ZTA): A Comprehensive Survey. *IEEE Access* **2022**, *10*, 57143–57179. https://doi.org/10.1109/access.2022.3174679.
26. Birkholz, H.; Thaler, D.; Richardson, M.; Smith, N.; Pan, W. Remote ATtestation procedureS (RATS) Architecture. Technical report, RFC Editor, 2023. https://doi.org/10.17487/rfc9334.
27. World Wide Web Consortium. PROV-DM: The PROV Data Model. W3C Recommendation, 2013. Accessed 6 August 2026.
28. World Wide Web Consortium. PROV-O: The PROV Ontology. W3C Recommendation, 2013. Accessed 6 August 2026.
29. Ahmed, A.I.A.; Ab Hamid, S.H.; Gani, A.; khan, S.; Khan, M.K. Trust and reputation for Internet of Things: Fundamentals, taxonomy, and open research challenges. *Journal of Network and Computer Applications* **2019**, *145*, 102409. https://doi.org/10.1016/j.jnca.2019.102409.
30. Aaqib, M.; Ali, A.; Chen, L.; Nibouche, O. IoT trust and reputation: a survey and taxonomy. *Journal of Cloud Computing* **2023**, *12*. https://doi.org/10.1186/s13677-023-00416-8.
31. Jøsang, A. *Subjective Logic*; Springer International Publishing, 2016. https://doi.org/10.1007/978-3-319-42337-1.
32. Nomikos, N.; Gkonis, P.K.; Bithas, P.S.; Trakadas, P. A Survey on UAV-Aided Maritime Communications: Deployment Considerations, Applications, and Future Challenges. *IEEE Open Journal of the Communications Society* **2023**, *4*, 56–78. https://doi.org/10.1109/ojcoms.2022.3225590.
33. Ali, M.F.; Jayakody, D.N.K.; Chursin, Y.A.; Affes, S.; Dmitry, S. Recent Advances and Future Directions on Underwater Wireless Communications. *Archives of Computational Methods in Engineering* **2020**, *27*, 1379–1412. https://doi.org/10.1007/s11831-019-09354-8.
34. Saeed, N.; Celik, A.; Al-Naffouri, T.Y.; Alouini, M.S. Underwater optical wireless communications, networking, and localization: A survey. *Ad Hoc Networks* **2019**, *94*, 101935. https://doi.org/10.1016/j.adhoc.2019.101935.
35. Jahanbakht, M.; Xiang, W.; Hanzo, L.; Rahimi Azghadi, M. Internet of Underwater Things and Big Marine Data Analytics—A Comprehensive Survey. *IEEE Communications Surveys & Tutorials* **2021**, *23*, 904–956. https://doi.org/10.1109/comst.2021.3053118.
36. Burleigh, S.; Fall, K.; Birrane, E.J. Bundle Protocol Version 7. RFC 9171, 2022. Accessed 22 October 2026.

37. Birrane, E.J.; McKeever, K. Bundle Protocol Security. RFC 9172, 2022. Accessed 22 October 2026.
38. Mui, L.; Mohtashemi, M.; Halberstadt, A. A computational model of trust and reputation. In Proceedings of the 35th Annual Hawaii International Conference on System Sciences. IEEE Comput. Soc, 2002, pp. 2431–2439. https://doi.org/10.1109/hicss.2002.994181.
39. Kamvar, S.D.; Schlosser, M.T.; Garcia-Molina, H. The Eigentrust algorithm for reputation management in P2P networks. In Proceedings of the twelfth international conference on World Wide Web - WWW '03. ACM Press, 2003, p. 640. https://doi.org/10.1145/775152.775242.
40. Wu, Y.; Low, K.H.; Lv, C. Cooperative Path Planning for Heterogeneous Unmanned Vehicles in a Search-and-Track Mission Aiming at an Underwater Target. *IEEE Transactions on Vehicular Technology* **2020**, *69*, 6782–6787. https://doi.org/10.1109/tvt.2020.2991983.
41. Ke, C.; Chen, H. Cooperative path planning for air–sea heterogeneous unmanned vehicles using search-and-tracking mission. *Ocean Engineering* **2022**, *262*, 112020. https://doi.org/10.1016/j.oceaneng.2022.112020.
42. Xue, K.; Wu, T. Distributed Consensus of USVs under Heterogeneous UAV-USV Multi-Agent Systems Cooperative Control Scheme. *Journal of Marine Science and Engineering* **2021**, *9*, 1314. https://doi.org/10.3390/jmse9111314.
43. Liu, M.; Li, Y.; Zhu, L.; Guo, Y.; Liu, B. Formation Control for Mixed-Order UAVs–USVs–UUVs Systems under Cooperative and Optimal Control. *Journal of Marine Science and Engineering* **2023**, *11*, 704. https://doi.org/10.3390/jmse11040704.
44. Yahuza, M.; Idris, M.Y.I.; Ahmedy, I.B.; Wahab, A.W.A.; Nandy, T.; Noor, N.M.; Bala, A. Internet of Drones Security and Privacy Issues: Taxonomy and Open Challenges. *IEEE Access* **2021**, *9*, 57243–57270. https://doi.org/10.1109/access.2021.3072030.
45. Adam, N.; Ali, M.; Naeem, F.; Ghazy, A.S.; Kaddoum, G. State-of-the-Art Security Schemes for the Internet of Underwater Things: A Holistic Survey. *IEEE Open Journal of the Communications Society* **2024**, *5*, 6561–6592. https://doi.org/10.1109/ojcoms.2024.3474290.
46. Jouhari, M.; Ibrahimi, K.; Tembine, H.; Ben-Othman, J. Underwater Wireless Sensor Networks: A Survey on Enabling Technologies, Localization Protocols, and Internet of Underwater Things. *IEEE Access* **2019**, *7*, 96879–96899. https://doi.org/10.1109/access.2019.2928876.
47. Mohsan, S.A.H.; Mazinani, A.; Othman, N.Q.H.; Amjad, H. Towards the internet of underwater things: a comprehensive survey. *Earth Science Informatics* **2022**, *15*, 735–764. https://doi.org/10.1007/s12145-021-00762-8.
48. Jiang, J.; Zhu, X.; Han, G.; Guizani, M.; Shu, L. A Dynamic Trust Evaluation and Update Mechanism Based on C4.5 Decision Tree in Underwater Wireless Sensor Networks. *IEEE Transactions on Vehicular Technology* **2020**, *69*, 9031–9040. https://doi.org/10.1109/tvt.2020.2999566.
49. Arifeen, M.M.; Bhakta, D.; Remu, S.R.H.; Islam, M.M.; Mahmud, M.; Kaiser, M.S. Hidden Markov Model based Trust Management Model for Underwater Wireless Sensor Networks. In Proceedings of the International Conference on Computing Advancements. ACM, 2020, pp. 1–5. https://doi.org/10.1145/3377049.3377054.
50. Junior, F.M.R.; Kamienski, C.A. A Survey on Trustworthiness for the Internet of Things. *IEEE Access* **2021**, *9*, 42493–42514. https://doi.org/10.1109/access.2021.3066457.
51. Liu, Y.; Wang, J.; Yan, Z.; Wan, Z.; Jäntti, R. A Survey on Blockchain-Based Trust Management for Internet of Things. *IEEE Internet of Things Journal* **2023**, *10*, 5898–5922. https://doi.org/10.1109/jiot.2023.3237893.
52. Shafique, A.; Mehmood, A.; Elhadef, M. Detecting Signal Spoofing Attack in UAVs Using Machine Learning Models. *IEEE Access* **2021**, *9*, 93803–93815. https://doi.org/10.1109/access.2021.3089847.
53. Dang, Y.; Benzaid, C.; Yang, B.; Taleb, T.; Shen, Y. Deep-Ensemble-Learning-Based GPS Spoofing Detection for Cellular-Connected UAVs. *IEEE Internet of Things Journal* **2022**, *9*, 25068–25085. https://doi.org/10.1109/jiot.2022.3195320.
54. Semanjski, S.; Semanjski, I.; De Wilde, W.; Muls, A. Use of Supervised Machine Learning for GNSS Signal Spoofing Detection with Validation on Real-World Meaconing and Spoofing Data—Part I. *Sensors* **2020**, *20*, 1171. https://doi.org/10.3390/s20041171.
55. Liu, W.; Yue, X.; Chen, Y.; Denoeux, T. Trusted Multi-View Deep Learning with Opinion Aggregation. *Proceedings of the AAAI Conference on Artificial Intelligence* **2022**, *36*, 7585–7593. https://doi.org/10.1609/aaai.v36i7.20724.
56. Singh, J.; Cobbe, J.; Norval, C. Decision Provenance: Harnessing Data Flow for Accountable Systems. *IEEE Access* **2019**, *7*, 6562–6574. https://doi.org/10.1109/access.2018.2887201.
57. Hardjono, T.; Smith, N. Towards an attestation architecture for blockchain networks. *World Wide Web* **2021**, *24*, 1587–1615. https://doi.org/10.1007/s11280-021-00869-4.
58. Figueiredo, S.; Silva, P.; Iacovazzi, A.; Holubenko, V.; Casal, J.; Calero, J.M.A.; Wang, Q.; Colarejo, P.; Armitt, R.L.; Inches, G.; et al. ARCADIAN-IoT - Enabling Autonomous Trust, Security and Privacy Management for IoT. In *Lecture Notes in Computer Science*; Springer International Publishing, 2023; pp. 348–359. https://doi.org/10.1007/978-3-031-20936-9_28.
59. Ott, S.; Kamhuber, M.; Pecholt, J.; Wessel, S. Universal Remote Attestation for Cloud and Edge Platforms. In Proceedings of the 18th International Conference on Availability, Reliability and Security. ACM, 2023, pp. 1–11. https://doi.org/10.1145/3600160.3600171.
60. Baciu, V.E.; Braeken, A.; Segers, L.; Silva, B.d. Secure Tiny Machine Learning on Edge Devices: A Lightweight Dual Attestation Mechanism for Machine Learning. *Future Internet* **2025**, *17*, 85. https://doi.org/10.3390/fi17020085.

61. Fagan, M.; Megas, K.N.; Scarfone, K.; Smith, M. IoT device cybersecurity capability core baseline. Technical report, National Institute of Standards and Technology, 2020. https://doi.org/10.6028/nist.ir.8259a.
62. Androulaki, E.; Barger, A.; Bortnikov, V.; Cachin, C.; Christidis, K.; De Caro, A.; Enyeart, D.; Ferris, C.; Laventman, G.; Manevich, Y.; et al. Hyperledger Fabric: A Distributed Operating System for Permissioned Blockchains. In Proceedings of the Thirteenth EuroSys Conference. ACM, 2018, pp. 1–15. https://doi.org/10.1145/3190508.3190538.
63. Hang, L.; Kim, D.H. Design and Implementation of an Integrated IoT Blockchain Platform for Sensing Data Integrity. *Sensors* **2019**, *19*, 2228. https://doi.org/10.3390/s19102228.
64. Honar Pajooh, H.; Rashid, M.A.; Alam, F.; Demidenko, S. IoT Big Data provenance scheme using blockchain on Hadoop ecosystem. *Journal of Big Data* **2021**, *8*. https://doi.org/10.1186/s40537-021-00505-y.
65. Liu, H.; Han, D.; Li, D. Fabric-iot: A Blockchain-Based Access Control System in IoT. *IEEE Access* **2020**, *8*, 18207–18218. https://doi.org/10.1109/access.2020.2968492.
66. Khan, P.W.; Byun, Y.C.; Park, N. IoT-Blockchain Enabled Optimized Provenance System for Food Industry 4.0 Using Advanced Deep Learning. *Sensors* **2020**, *20*, 2990. https://doi.org/10.3390/s20102990.
67. Rahman, M.A.; Hossain, M.S.; Islam, M.S.; Alrajeh, N.A.; Muhammad, G. Secure and Provenance Enhanced Internet of Health Things Framework: A Blockchain Managed Federated Learning Approach. *IEEE Access* **2020**, *8*, 205071–205087. https://doi.org/10.1109/access.2020.3037474.
68. Suhail, S.; Hussain, R.; Jurdak, R.; Hong, C.S. Trustworthy Digital Twins in the Industrial Internet of Things With Blockchain. *IEEE Internet Computing* **2022**, *26*, 58–67. https://doi.org/10.1109/mic.2021.3059320.
69. Laurie, B.; Langley, A.; Kasper, E. Certificate Transparency. RFC 6962, 2013. Accessed 22 October 2026.
70. Laurie, B.; Messeri, E.; Stradling, R. Certificate Transparency Version 2.0. RFC 9162, 2021. Accessed 22 October 2026.
71. Newman, Z.; Meyers, J.S.; Torres-Arias, S. Sigstore: Software Signing for Everybody. In Proceedings of the 2022 ACM SIGSAC Conference on Computer and Communications Security, 2022, pp. 2353–2367. https://doi.org/10.1145/3548606.3560596.
72. Kephart, J.; Chess, D. The vision of autonomic computing. *Computer* **2003**, *36*, 41–50. https://doi.org/10.1109/MC.2003.1160055.
73. IBM. An Architectural Blueprint for Autonomic Computing, Third Edition, 2005. Accessed 22 October 2026.
74. Steinberg, A.N.; Bowman, C.L.; White, F.E. Revisions to the JDL data fusion model. In Proceedings of the SPIE Proceedings, 1999, Vol. 3719, p. 430. https://doi.org/10.1117/12.341367.
75. Wust, K.; Gervais, A. Do you Need a Blockchain? In Proceedings of the 2018 Crypto Valley Conference on Blockchain Technology (CVCBT), 2018, pp. 45–54. https://doi.org/10.1109/CVCBT.2018.00011.
76. Locher, T.; Obermeier, S.; Pignolet, Y.A. When Can a Distributed Ledger Replace a Trusted Third Party? In Proceedings of the 2018 IEEE International Conference on Internet of Things (iThings) and IEEE Green Computing and Communications (GreenCom) and IEEE Cyber, Physical and Social Computing (CPSCom) and IEEE Smart Data (SmartData), 2018, pp. 1069–1077. https://doi.org/10.1109/Cybermatics_2018.2018.00197.
77. Maersk. Maersk and IBM to Discontinue TradeLens, a Blockchain-Enabled Global Trade Platform, 2022. Accessed 22 October 2026.
78. Švarcmajer, M.; Kohler, M.; Krpić, Z.; Lukić, I. Consensus-Driven Framework for Data-Driven Optimization of Distributed Systems Through Blockchain Consensus Mechanism Selection. *Big Data and Cognitive Computing* **2026**, *10*, 154. https://doi.org/10.3390/bdcc10050154.
79. Ongaro, D.; Ousterhout, J. In Search of an Understandable Consensus Algorithm. In Proceedings of the 2014 USENIX Annual Technical Conference, 2014, pp. 305–319.
80. Castro, M.; Liskov, B. Practical Byzantine Fault Tolerance. In Proceedings of the Third Symposium on Operating Systems Design and Implementation, 1999, pp. 173–186.
81. Yin, M.; Malkhi, D.; Reiter, M.K.; Gueta, G.G.; Abraham, I. HotStuff: BFT Consensus with Linearity and Responsiveness. In Proceedings of the 2019 ACM Symposium on Principles of Distributed Computing, 2019, pp. 347–356. https://doi.org/10.1145/3293611.3331591.
82. CometBFT Contributors. Consensus Specification, 2026. Accessed 22 October 2026.
83. Danezis, G.; Kokoris-Kogias, L.; Sonnino, A.; Spiegelman, A. Narwhal and Tusk: a DAG-based mempool and efficient BFT consensus. In Proceedings of the Seventeenth European Conference on Computer Systems, 2022, pp. 34–50. https://doi.org/10.1145/3492321.3519594.
84. Dorri, A.; Kanhere, S.S.; Jurdak, R. Towards an Optimized BlockChain for IoT. In Proceedings of the Second International Conference on Internet-of-Things Design and Implementation. ACM, 2017, pp. 173–178. https://doi.org/10.1145/3054977.3055003.
85. Gilbert, S.; Lynch, N. Brewer's Conjecture and the Feasibility of Consistent, Available, Partition-Tolerant Web Services. *ACM SIGACT News* **2002**, *33*, 51–59. https://doi.org/10.1145/564585.564601.
86. Hyperledger Fabric Contributors. Configuring and Operating a BFT Ordering Service, 2026. Accessed 22 October 2026.
87. Hyperledger Foundation. Hyperledger Fabric Documentation: Ordering Service, 2026. Accessed 6 August 2026.

88. Josefsson, S.; Liusvaara, I. Edwards-Curve Digital Signature Algorithm (EdDSA). RFC 8032, 2017. Accessed 22 October 2026.
89. National Institute of Standards and Technology. Module-Lattice-Based Digital Signature Standard. FIPS 204, 2024. Accessed 22 October 2026.
90. Neto, E.C.P.; Dadkhah, S.; Ferreira, R.; Zohourian, A.; Lu, R.; Ghorbani, A.A. CICIoT2023: A Real-Time Dataset and Benchmark for Large-Scale Attacks in IoT Environment. *Sensors* **2023**, *23*, 5941. https://doi.org/10.3390/s23135941.
91. Alsaedi, A.; Moustafa, N.; Tari, Z.; Mahmood, A.; Anwar, A. TON_IoT Telemetry Dataset: A New Generation Dataset of IoT and IIoT for Data-Driven Intrusion Detection Systems. *IEEE Access* **2020**, *8*, 165130–165150. https://doi.org/10.1109/access.2020.3022862.
92. Ferrag, M.A.; Friha, O.; Hamouda, D.; Maglaras, L.; Janicke, H. Edge-IIoTset: A New Comprehensive Realistic Cyber Security Dataset of IoT and IIoT Applications for Centralized and Federated Learning. *IEEE Access* **2022**, *10*, 40281–40306. https://doi.org/10.1109/access.2022.3165809.
93. Varga, L.A.; Kiefer, B.; Messmer, M.; Zell, A. SeaDronesSee: A Maritime Benchmark for Detecting Humans in Open Water. In Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision, 2022, pp. 2260–2270.
94. Shah, S.; Dey, D.; Lovett, C.; Kapoor, A. AirSim: High-Fidelity Visual and Physical Simulation for Autonomous Vehicles. In *Springer Proceedings in Advanced Robotics*; Springer International Publishing, 2018; pp. 621–635. https://doi.org/10.1007/978-3-319-67361-5_40.
95. Manhaes, M.M.M.; Scherer, S.A.; Voss, M.; Douat, L.R.; Rauschenbach, T. UUV Simulator: A Gazebo-based package for underwater intervention and multi-robot simulation. In Proceedings of the OCEANS 2016 MTS/IEEE Monterey. IEEE, 2016, pp. 1–8. https://doi.org/10.1109/oceans.2016.7761080.
96. Martin, R.; Rajasekaran, S.; Peng, Z. Aqua-Sim Next Generation. In Proceedings of the International Conference on Underwater Networks & Systems. ACM, 2017, pp. 1–8. https://doi.org/10.1145/3148675.3148679.
97. Zhou, Y.; Soh, Y.S.; Loh, H.S.; Yuen, K.F. The key challenges and critical success factors of blockchain implementation: Policy implications for Singapore's maritime industry. *Marine Policy* **2020**, *122*, 104265. https://doi.org/10.1016/j.marpol.2020.104265.
98. Kshetri, N. Blockchain's roles in meeting key supply chain management objectives. *International Journal of Information Management* **2018**, *39*, 80–89. https://doi.org/10.1016/j.ijinfomgt.2017.12.005.
99. Ølnes, S.; Ubacht, J.; Janssen, M. Blockchain in government: Benefits and implications of distributed ledger technology for information sharing. *Government Information Quarterly* **2017**, *34*, 355–364. https://doi.org/10.1016/j.giq.2017.09.007.
100. Reyna, A.; Martín, C.; Chen, J.; Soler, E.; Díaz, M. On blockchain and its integration with IoT. Challenges and opportunities. *Future Generation Computer Systems* **2018**, *88*, 173–190. https://doi.org/10.1016/j.future.2018.05.046.
101. Kshetri, N. Can Blockchain Strengthen the Internet of Things? *IT Professional* **2017**, *19*, 68–72. https://doi.org/10.1109/mitp.2017.3051335.
102. Safety-Critical Systems Club. GSN Community Standard, Version 3. SCSC-141C, 2021. Accessed 22 October 2026.
103. Object Management Group. Structured Assurance Case Metamodel (SACM), Version 2.3, 2023. Accessed 22 October 2026.
104. UL Standards and Engagement. UL 4600: Evaluation of Autonomous Products, Third Edition, 2023. Accessed 22 October 2026.
105. International Maritime Organization. Interim Guidelines for MASS Trials. MSC.1/Circ.1604, 2019. Accessed 22 October 2026.
106. International Maritime Organization. Outcome of the Regulatory Scoping Exercise for the Use of Maritime Autonomous Surface Ships (MASS). MSC.1/Circ.1638, 2021. Accessed 22 October 2026.
107. International Maritime Organization. International Code of Safety for Maritime Autonomous Surface Ships (MASS Code). Resolution MSC.595(111), 2026. Accessed 22 October 2026.
108. DNV. Autonomous and Remotely Operated Ships. DNV-CG-0264, 2024. Accessed 22 October 2026.
109. Moran, B.; Tschofenig, H.; Brown, D.; Meriac, M. A Firmware Update Architecture for Internet of Things. RFC 9019, 2021. https://doi.org/10.17487/RFC9019.
110. 3rd Generation Partnership Project. Unmanned Aerial System (UAS) Support in 3GPP; Stage 1. 3GPP TS 22.125, 2026. Specification record accessed 6 August 2026.
111. 3rd Generation Partnership Project. Maritime Communication Services over 3GPP System. 3GPP TS 22.119, 2026. Specification record accessed 6 August 2026.